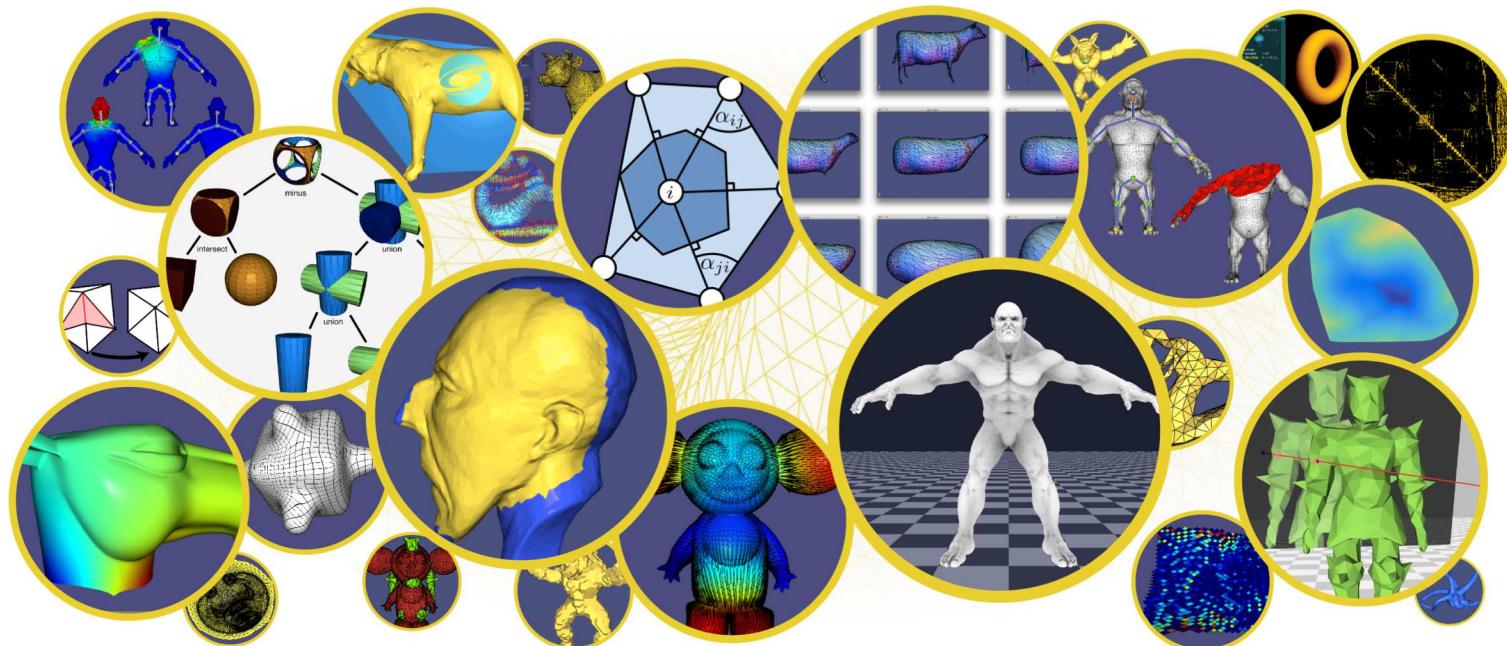


libigl Prototyping Geometry Processing Research in C++



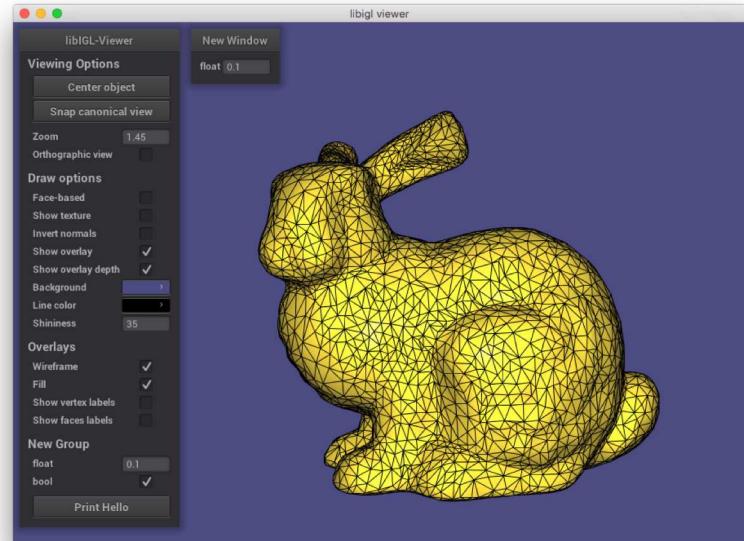
Alec Jacobson
University of Toronto

<https://github.com/libigl/libigl-course>

Daniele Panozzo
New York University

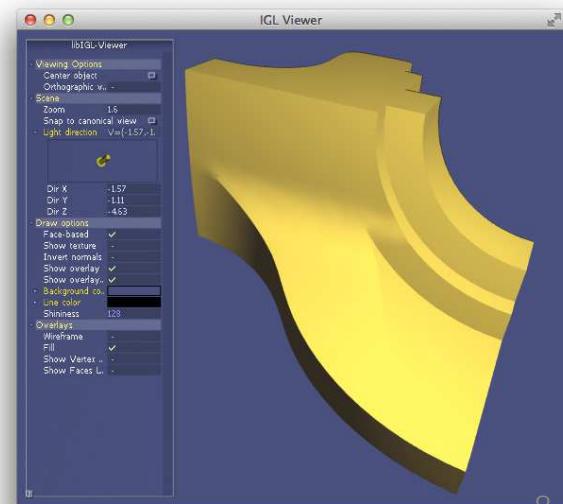
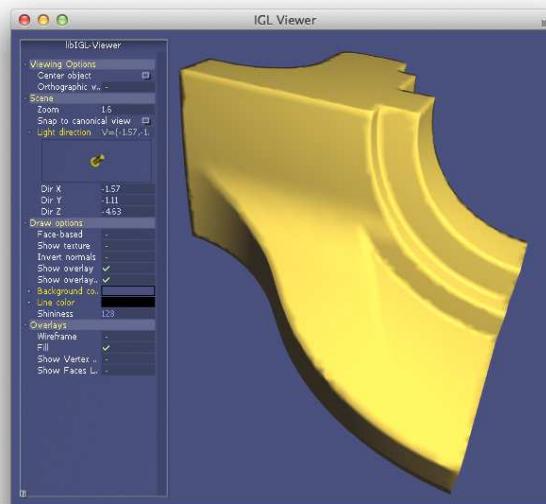
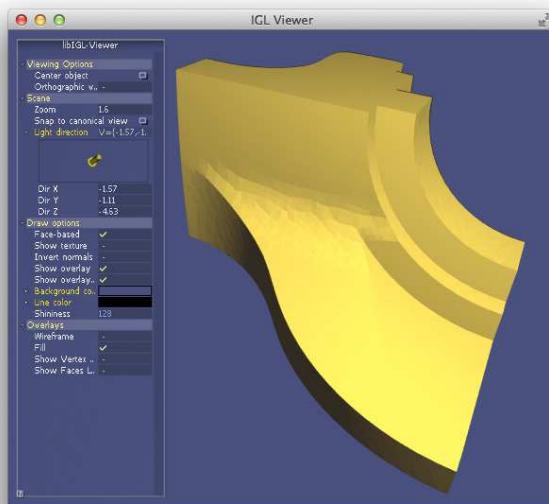
libigl is amassing a large library of C++ implementations of core routines

- mesh I/O (.obj, .stl, .off, .ply, .wrl, .mesh)



libigl is amassing a large library of C++ implementations of core routines

- mesh I/O (.obj, .stl, .off, .ply, .wrl, .mesh)
- areas, normals, angles, edge lengths, adjacencies



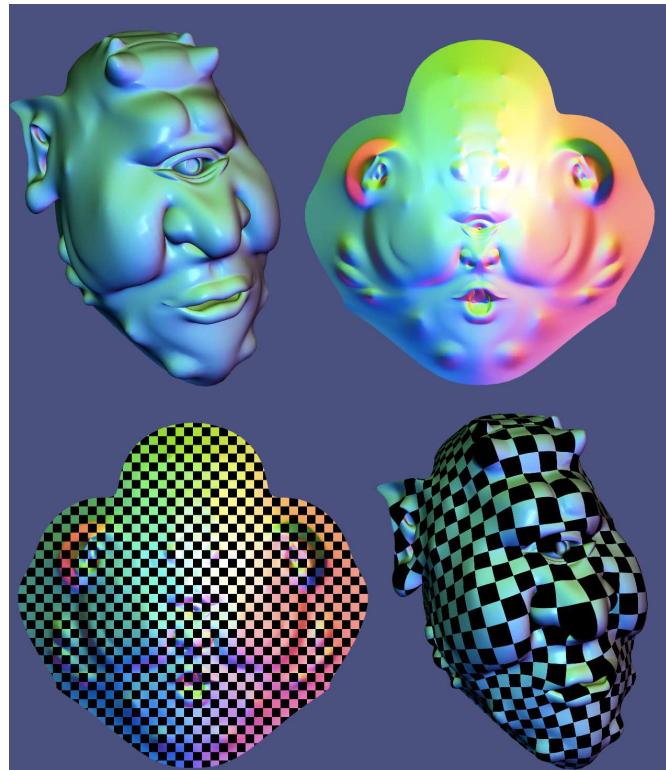
libigl is amassing a large library of C++ implementations of core routines

- mesh I/O (.obj, .stl, .off, .ply, .wrl, .mesh)
- areas, normals, angles, edge lengths, adjacencies
- cotangent Laplacian, mass matrix, gradient, divergence



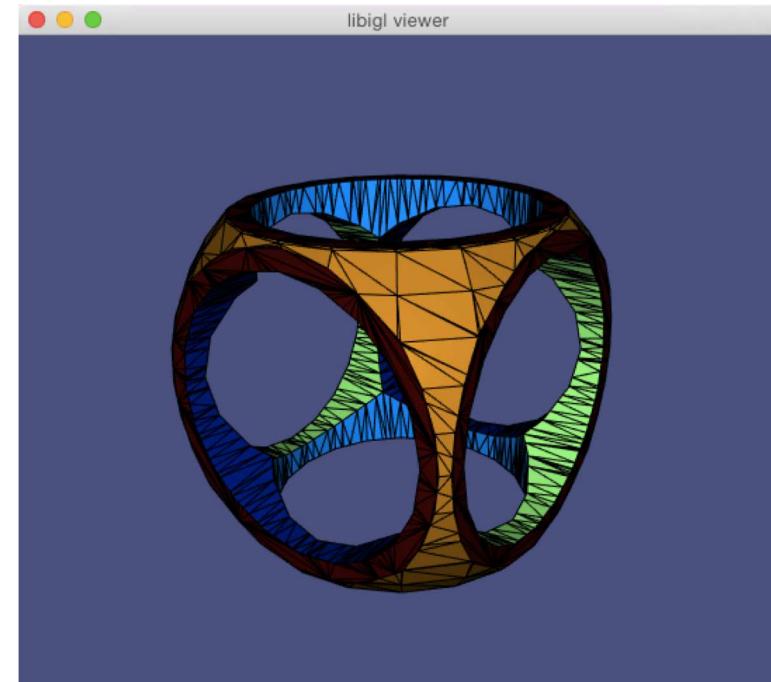
libigl is amassing a large library of C++ implementations of core routines

- mesh I/O (.obj, .stl, .off, .ply, .wrl, .mesh)
- areas, normals, angles, edge lengths, adjacencies
- cotangent Laplacian, mass matrix, gradient, divergence
- parameterization, smoothing, deformation, decimation



libigl is amassing a large library of C++ implementations of core routines

- mesh I/O (.obj, .stl, .off, .ply, .wrl, .mesh)
- areas, normals, angles, edge lengths, adjacencies
- cotangent Laplacian, mass matrix, gradient, divergence
- parameterization, smoothing, deformation, decimation
- boolean operations, quad meshing, closest point queries



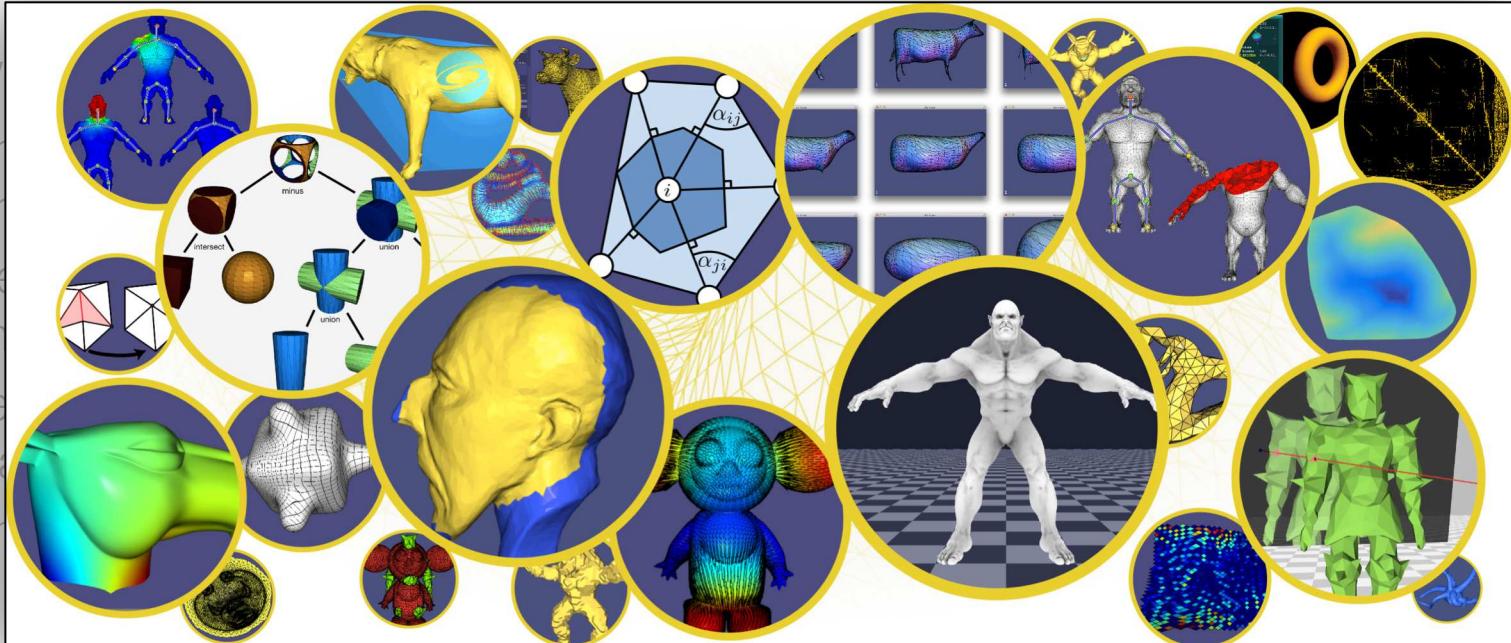
libigl is amassing a large library of C++ implementations of core routines

- mesh I/O (.obj, .stl, .off, .ply, .wrl, .mesh)
- areas, normals, angles, edge lengths, adjacencies
- cotangent Laplacian, mass matrix, gradient, divergence
- parameterization, smoothing, deformation, decimation
- boolean operations, quad meshing, closest point queries
- CGAL, tetgen, python, matlab



libigl is amassing a large library of C++ implementations of core routines

- mesh I/O
- areas, normals, adjacencies
- cotangent, gradient
- parameterizations
- deformations
- boolean operations
- closest points
- CGAL, tetgen, python, matlab
- ... much more



libigl continues its mission to empower *research*



libigl continues its mission to empower *research*



UNIVERSITY OF
TORONTO



Berkeley
UNIVERSITY OF CALIFORNIA



ETH zürich



UNIVERSITÉ
TOULOUSE III
PAUL SABATIER



THE HONG KONG
UNIVERSITY OF SCIENCE
AND TECHNOLOGY



UCL



TEXAS
The University of Texas at Austin



Emphasis on fast prototyping and exploration, rather than app development

libigl hosts an online tutorial

libigl.github.io/libigl/tutorial/tutorial.html

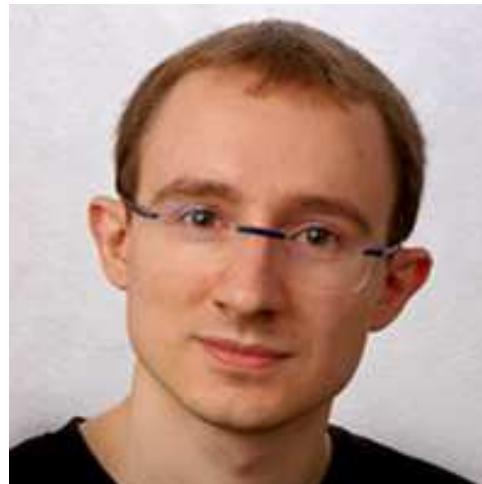


Libigl is an open source C++ library for geometry processing research and development. Dropping the heavy data structures of tradition geometry libraries, libigl is a simple header-only library of encapsulated functions. This combines

Who are we?

Alec Jacobson

University of Toronto



Daniele Panozzo

New York University

James Zhou

Adobe Research



Winding Number
Mesh Booleans
Unit Testing

Jeremie Dumas

NYU/nTopology



CMake Build System
Viewer

Sebastian Koch

TU Berlin

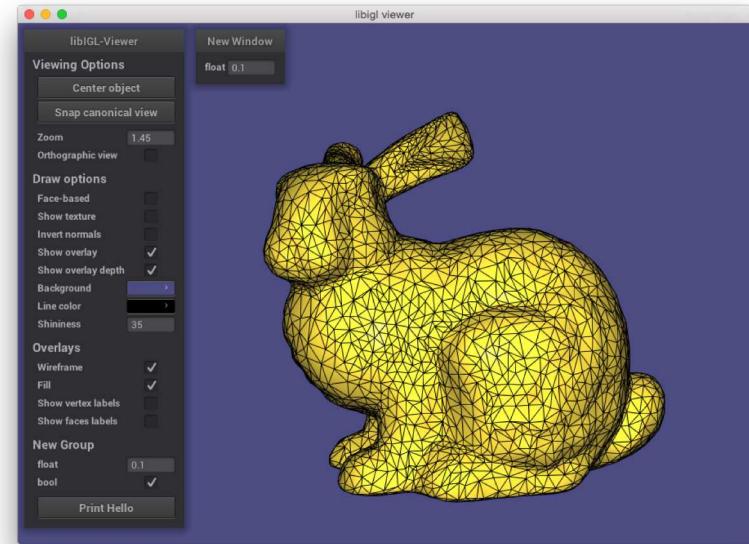
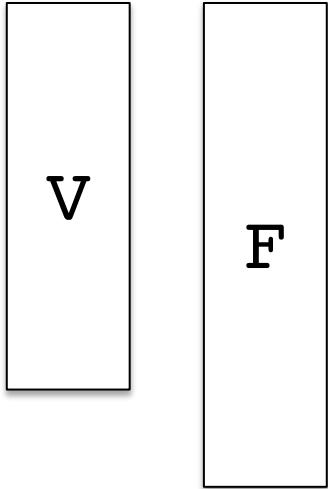


Python Bindings
Javascript Port
(coming soon!)

Community Effort



Christian Schüller,
Kevin Wallimann,
Stefan Brugger,
Yotam Gingold,
....

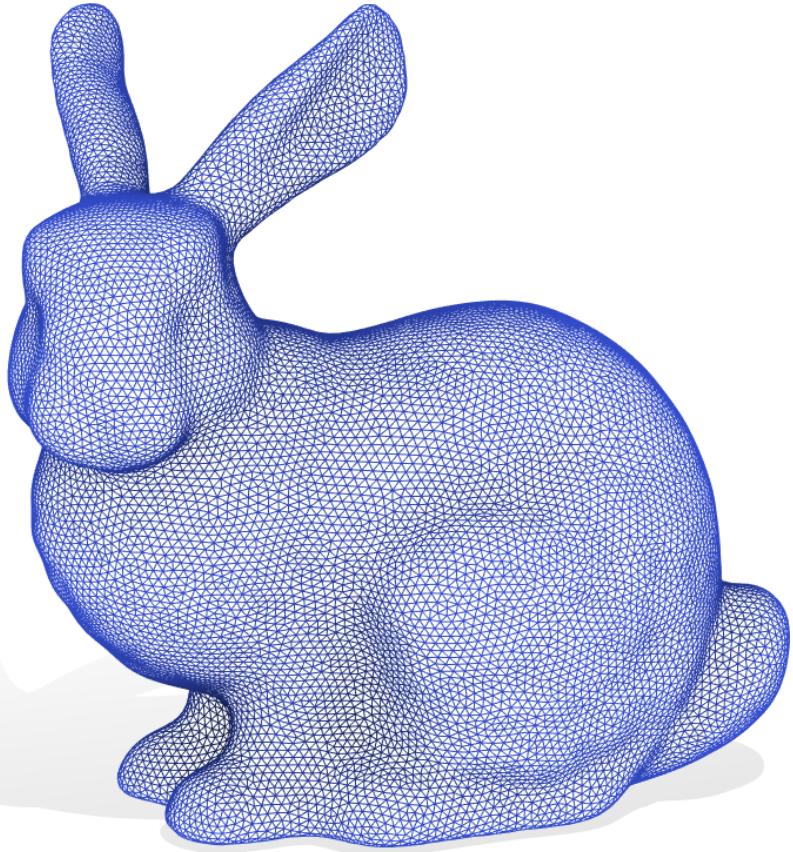


libigl data structures & style

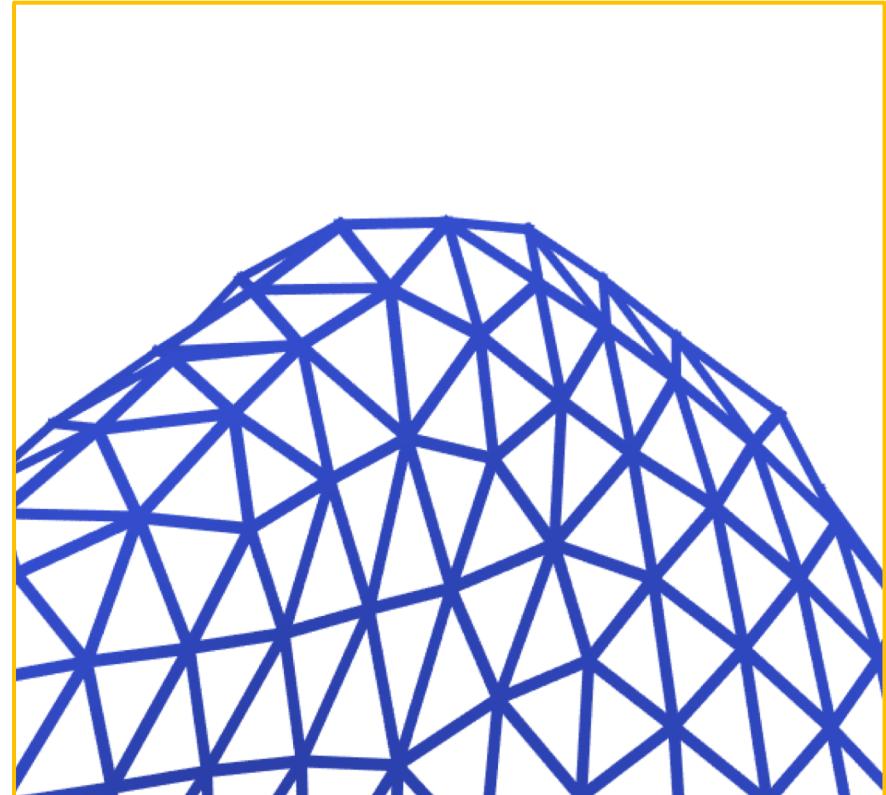
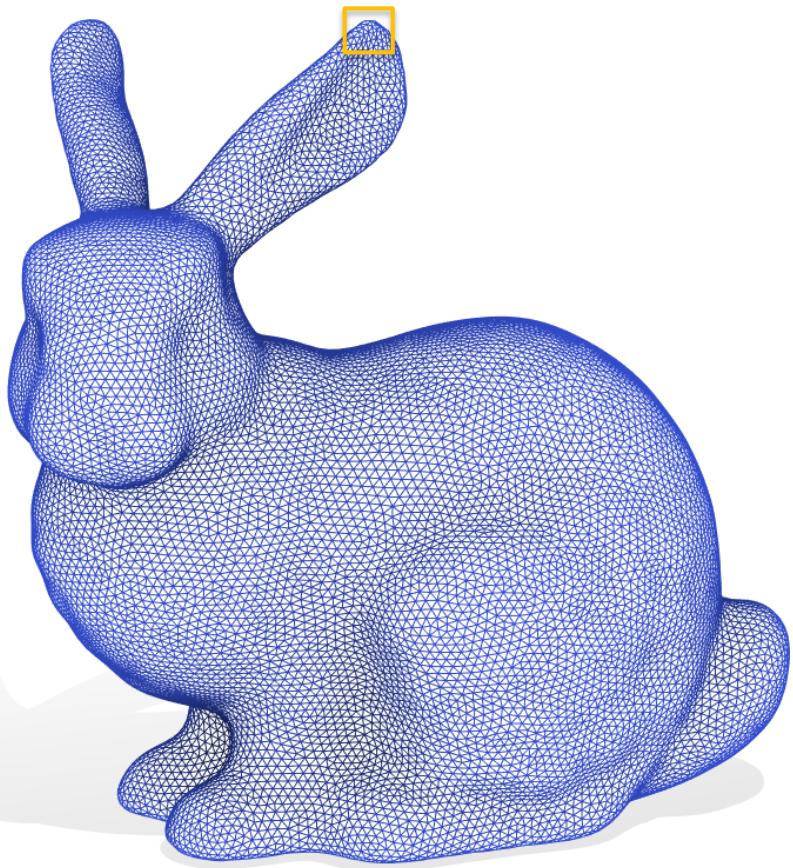
Triangle meshes discretize surfaces...



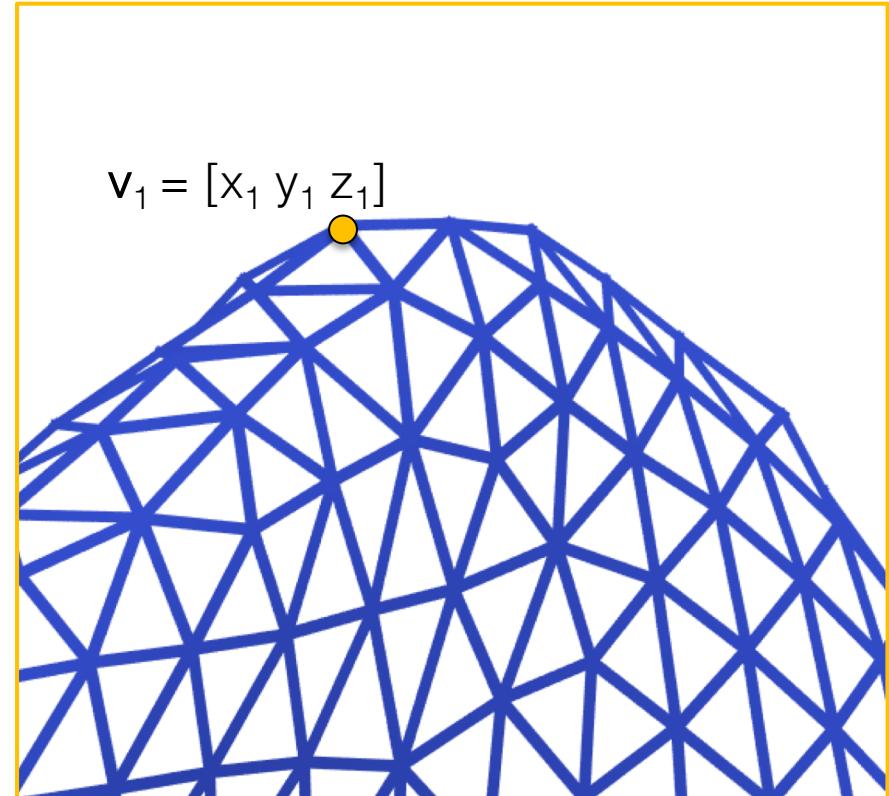
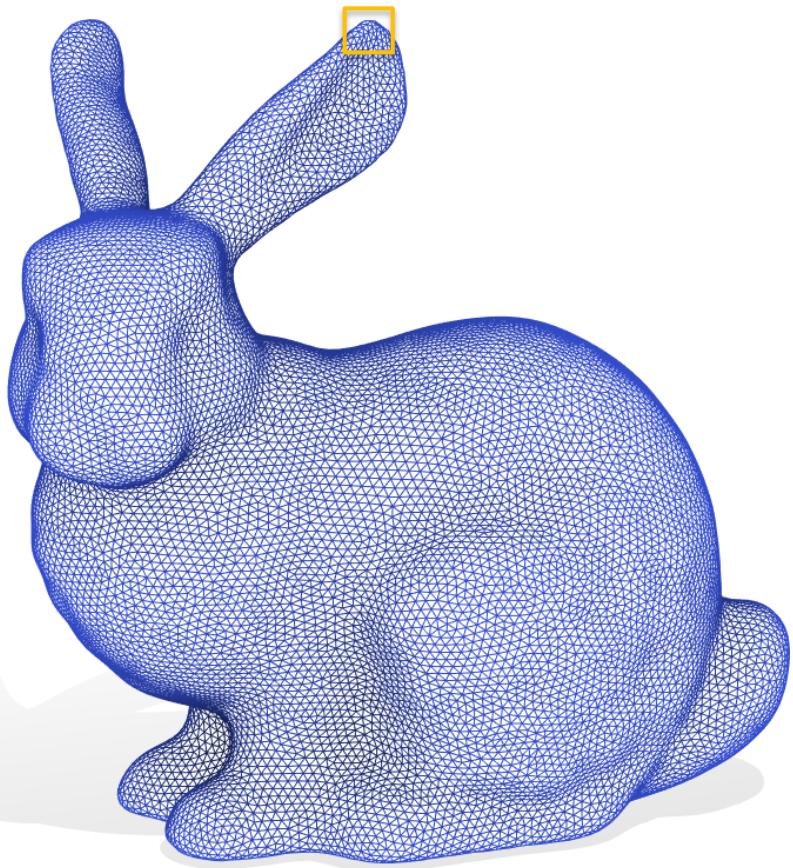
Triangle meshes discretize surfaces...



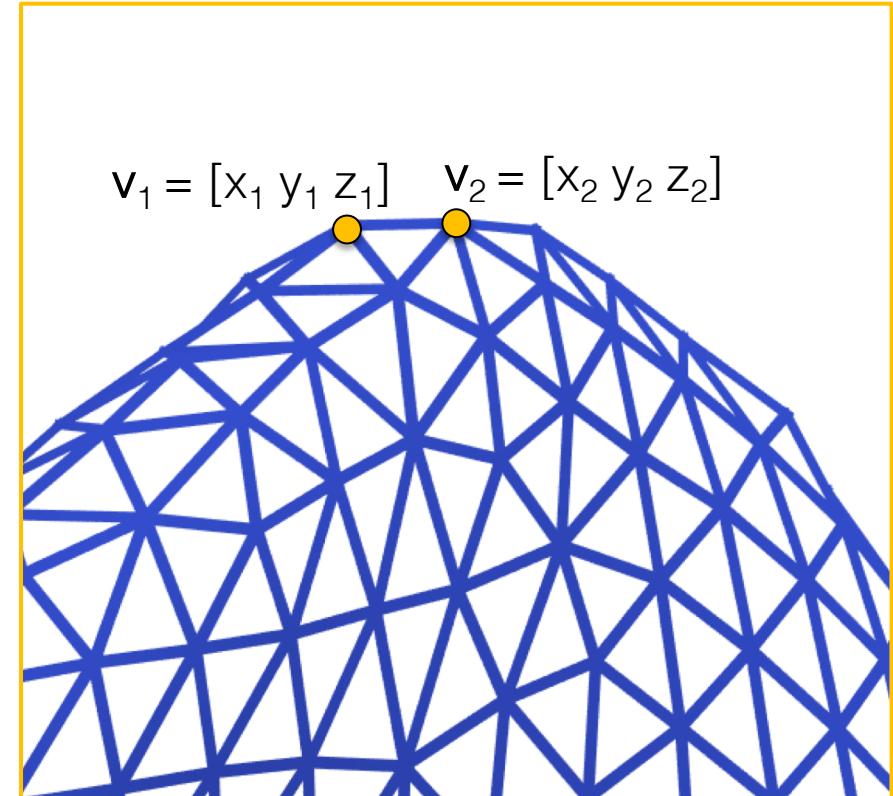
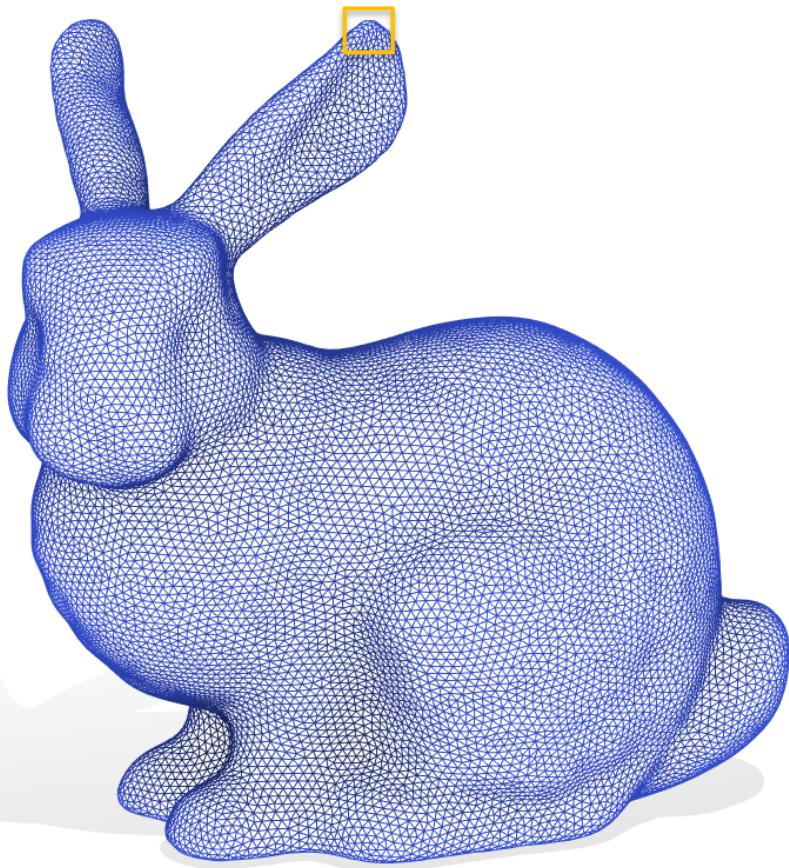
Triangle meshes discretize surfaces...



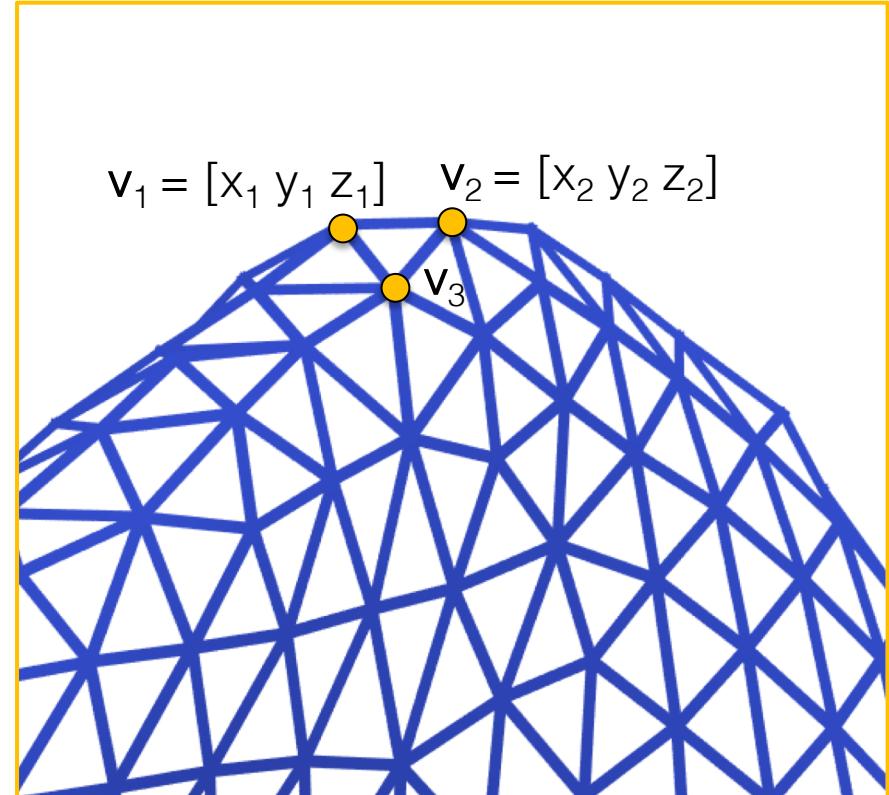
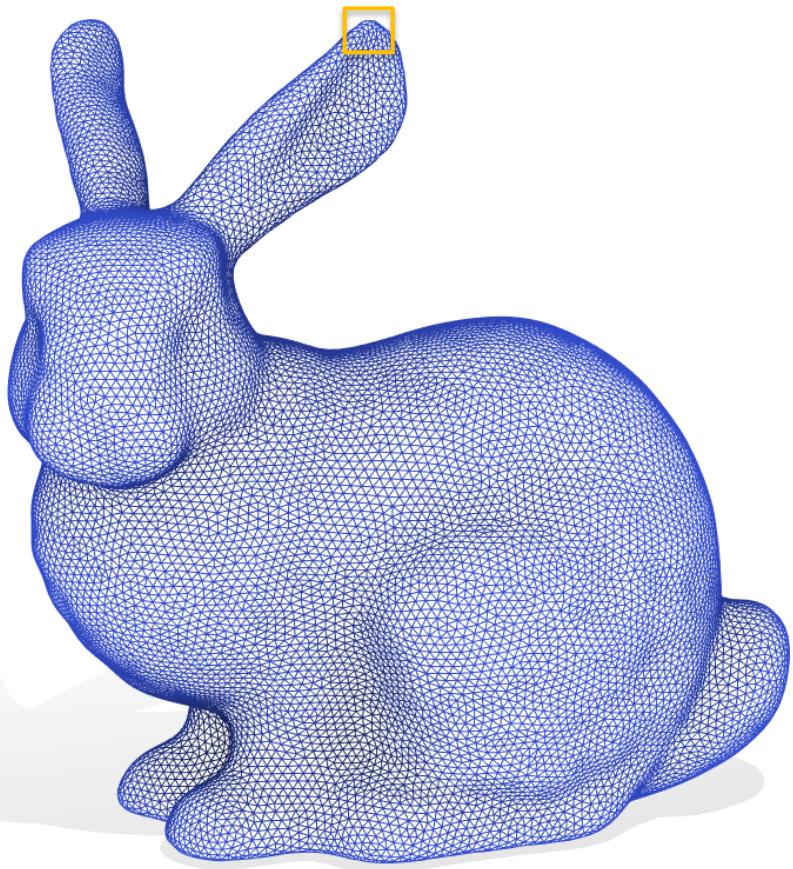
Triangle meshes discretize surfaces...



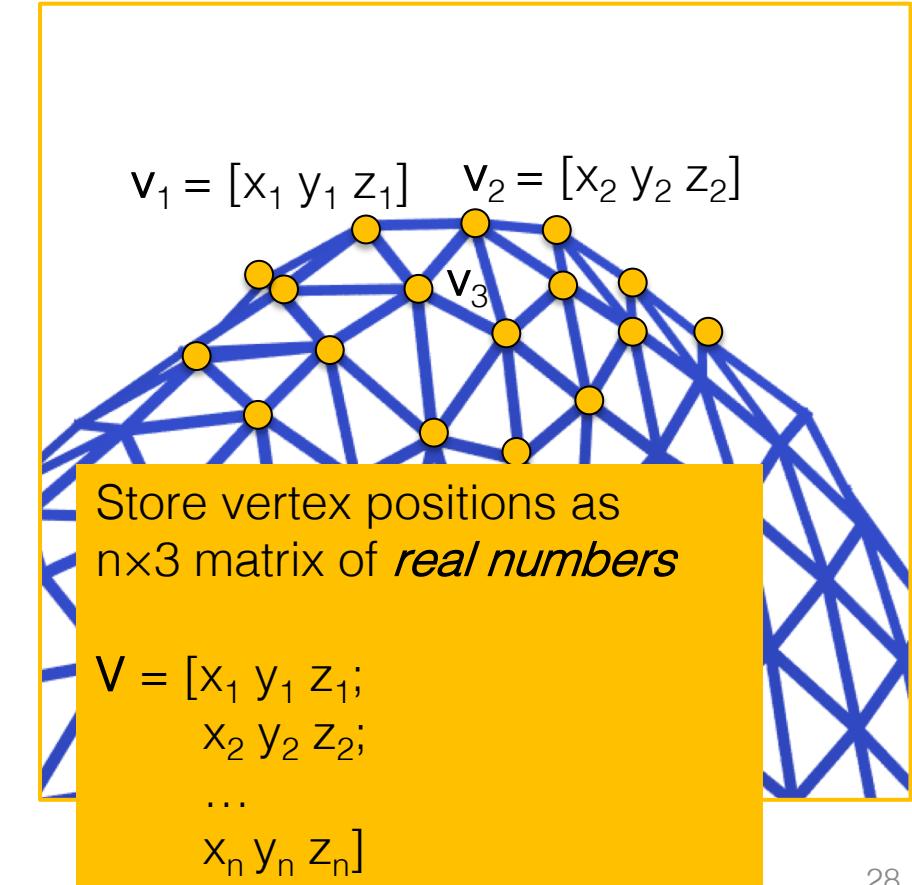
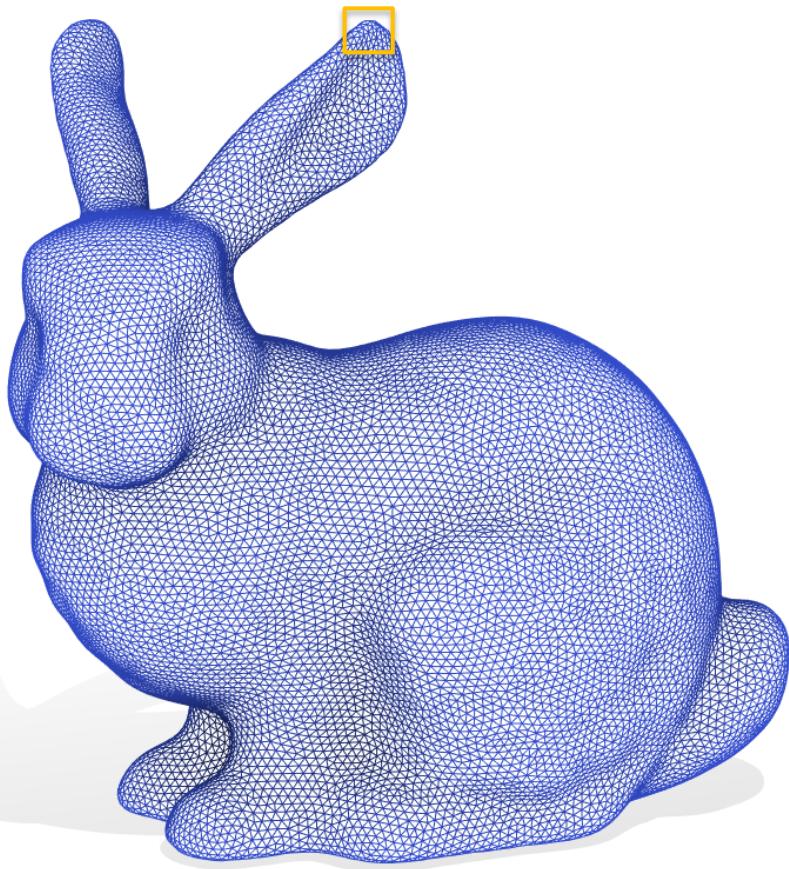
Triangle meshes discretize surfaces...



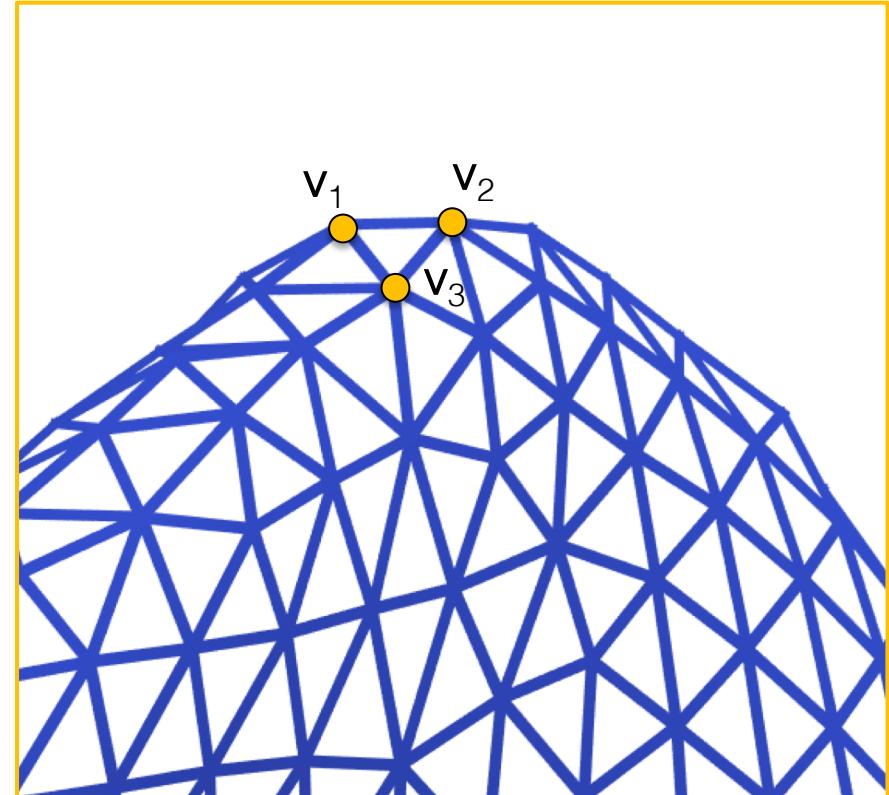
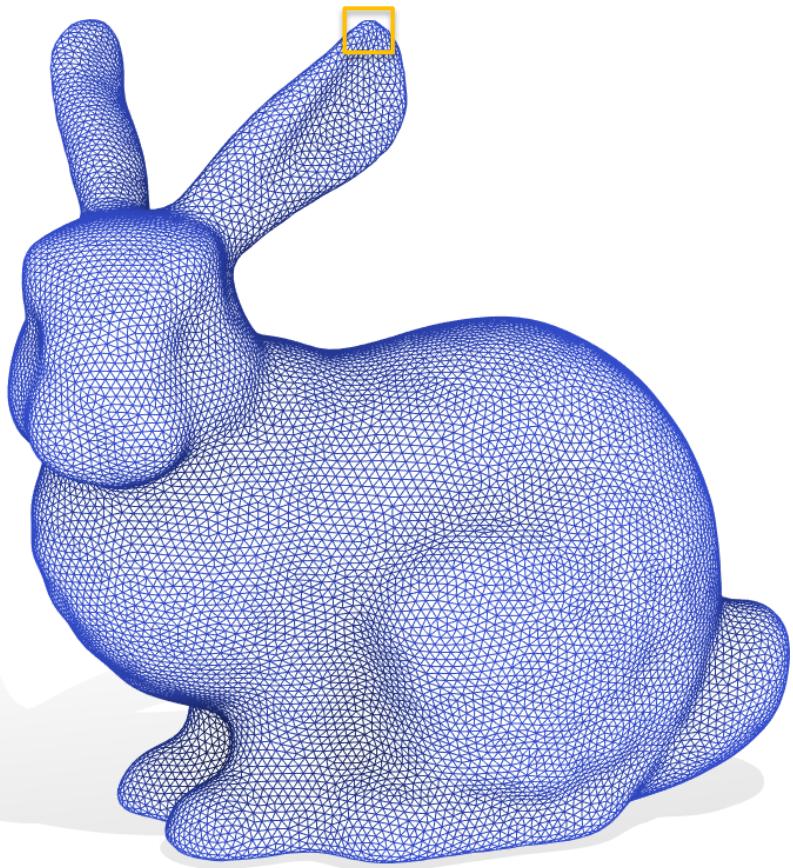
Triangle meshes discretize surfaces...



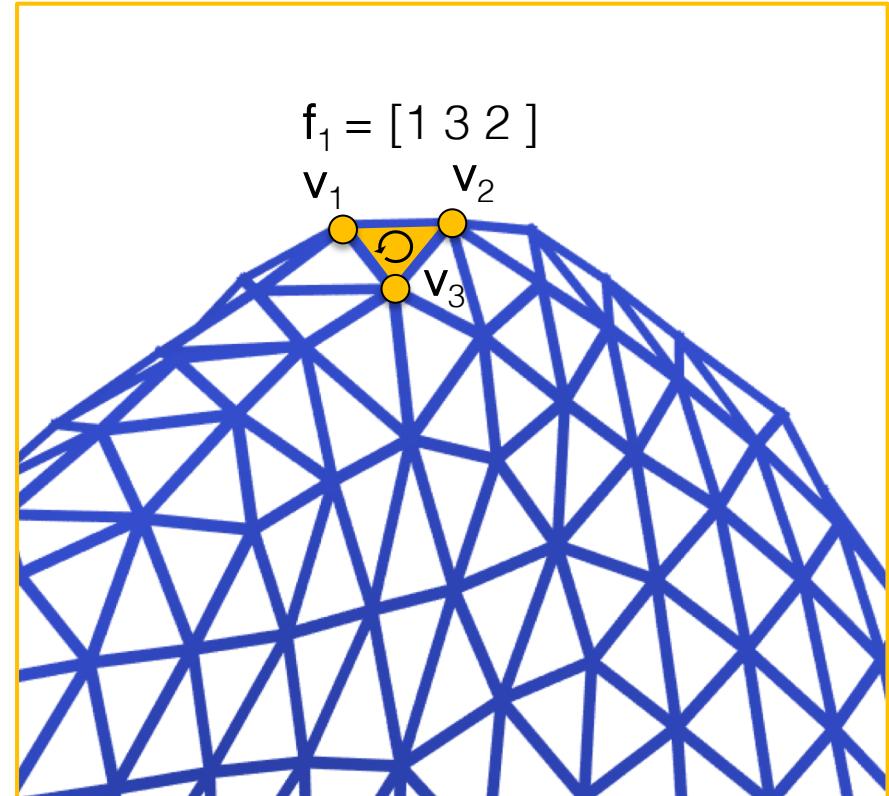
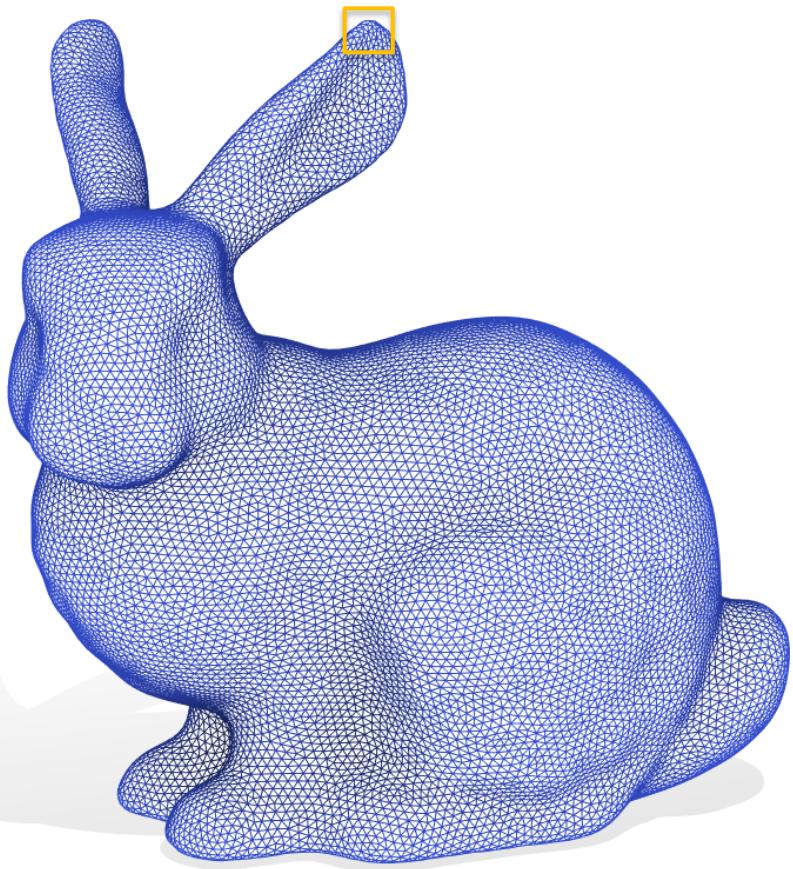
Triangle meshes discretize surfaces...



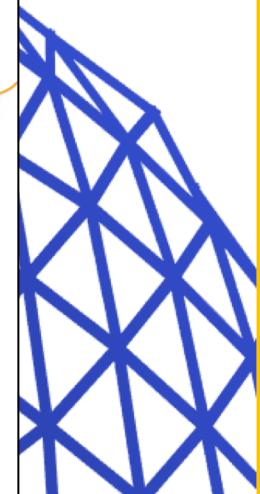
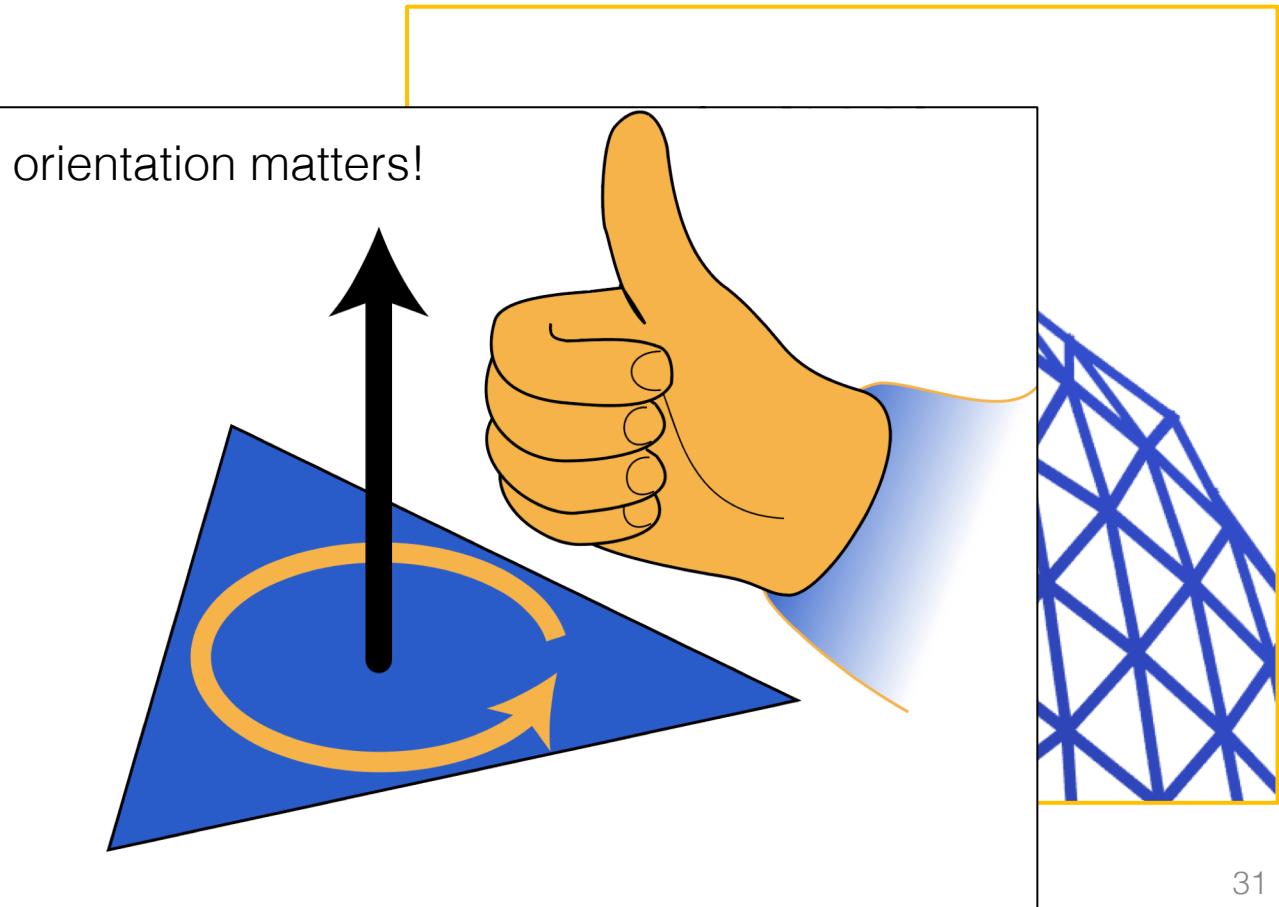
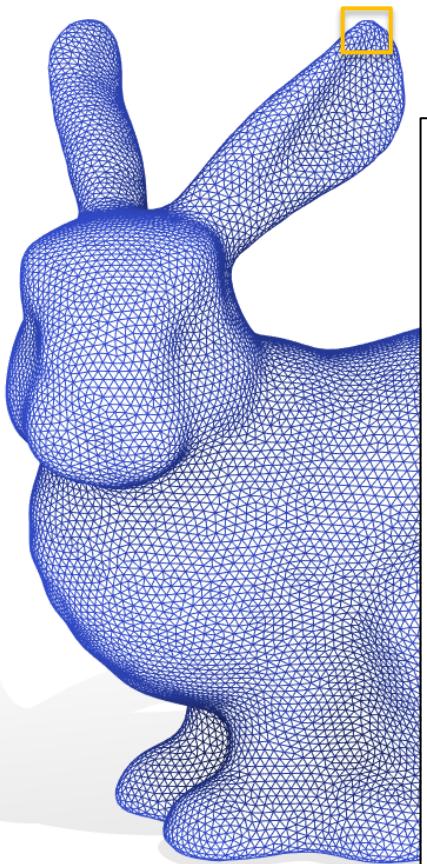
Triangle meshes discretize surfaces...



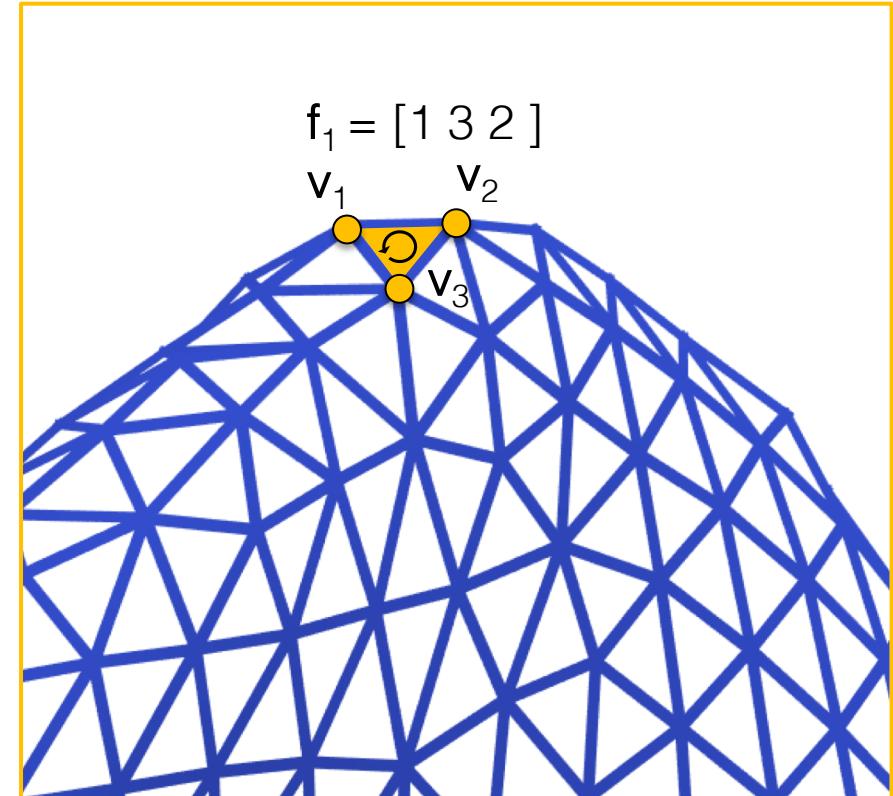
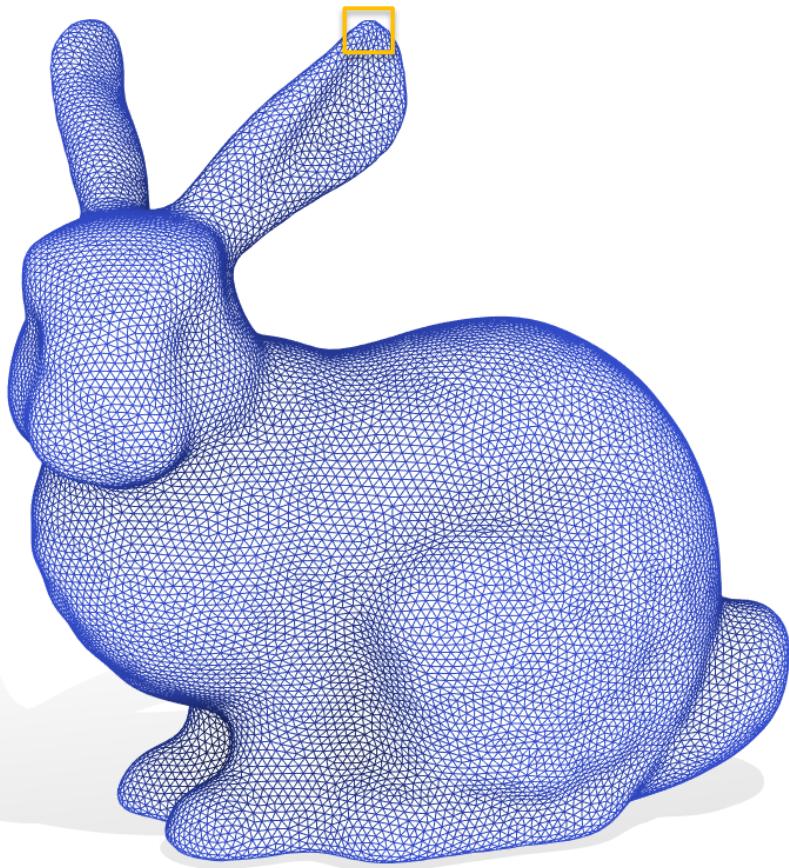
Triangle meshes discretize surfaces...



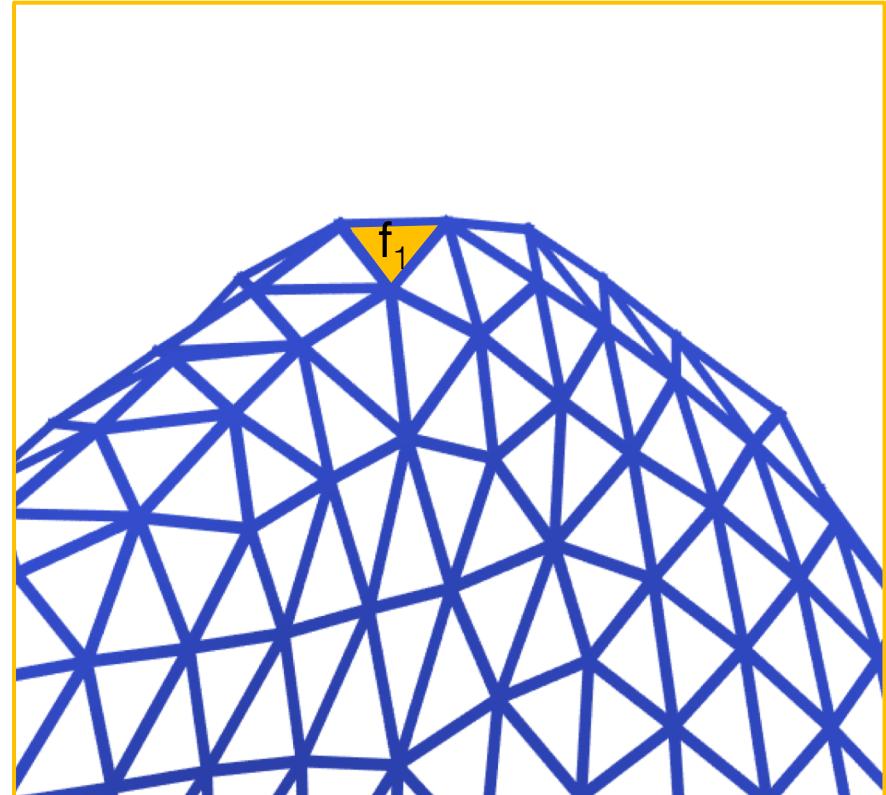
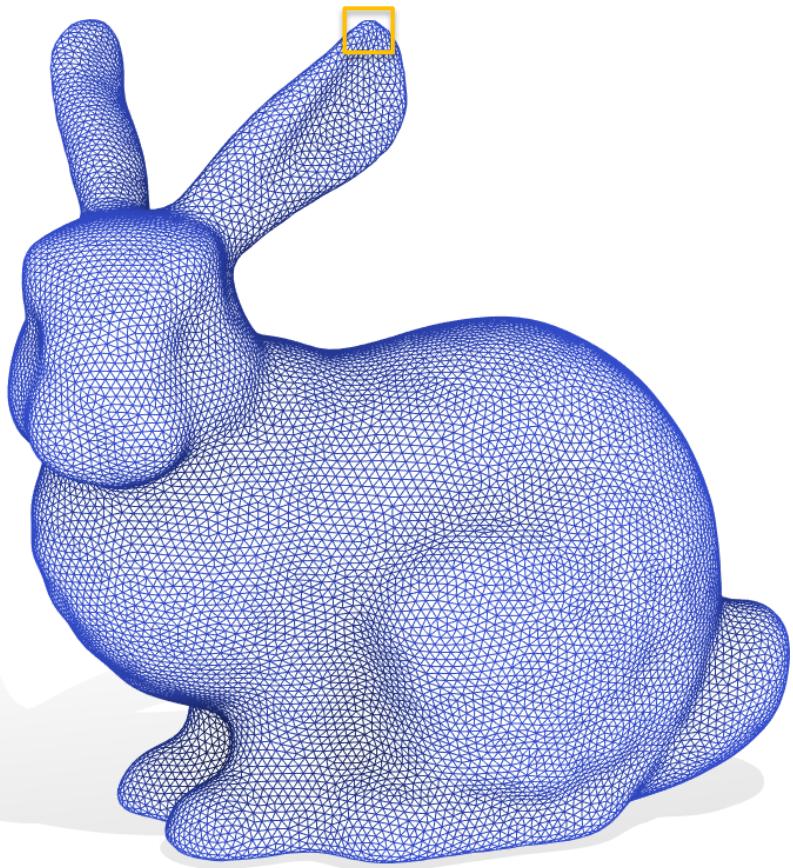
Triangle meshes discretize surfaces...



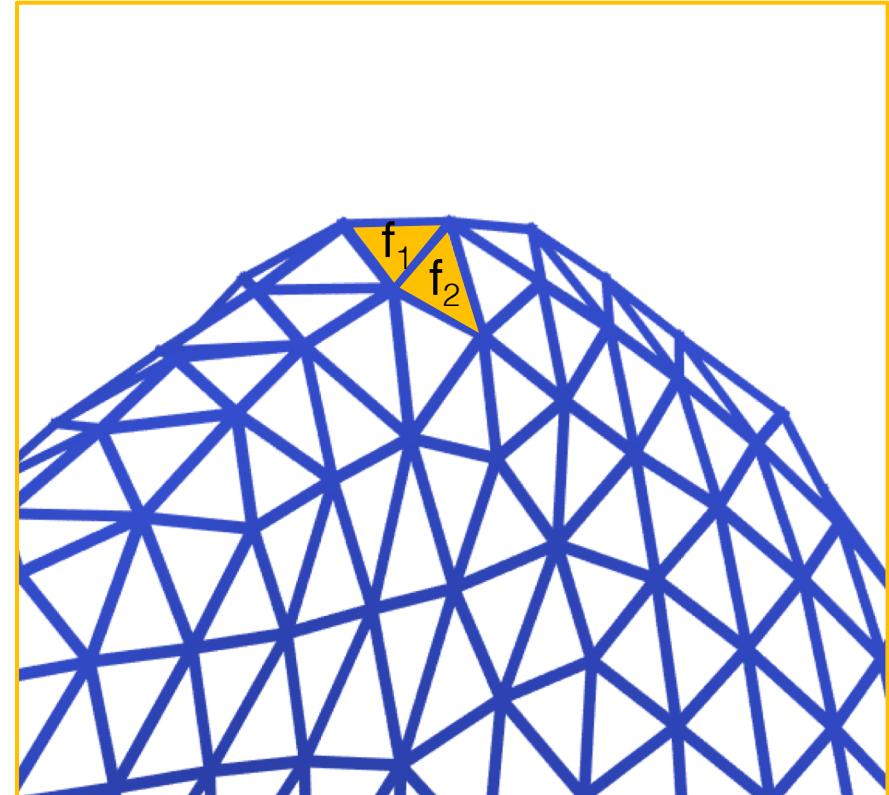
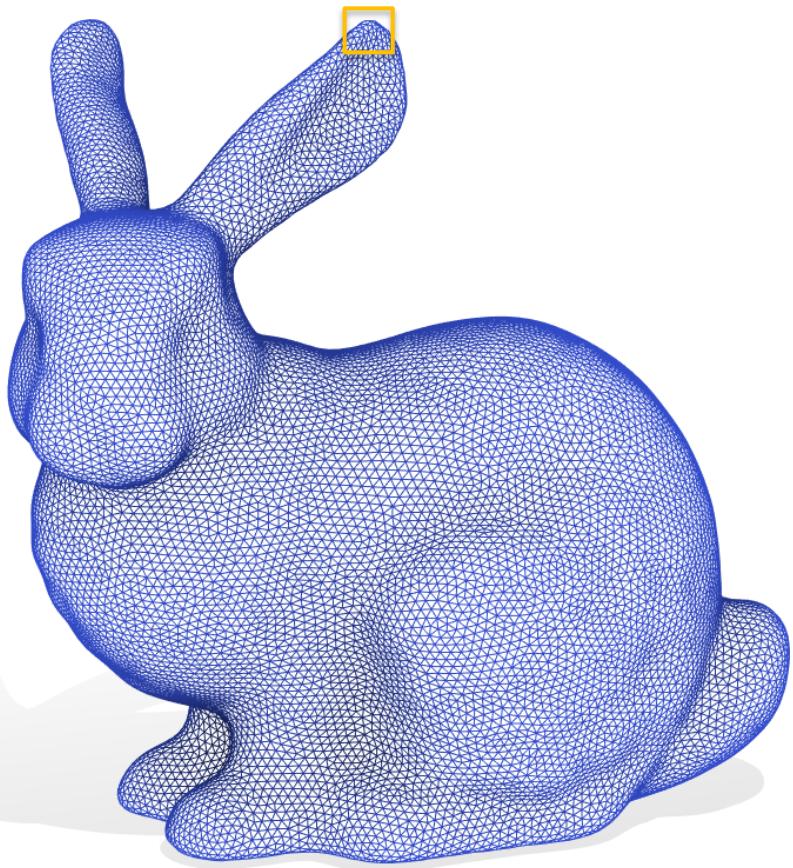
Triangle meshes discretize surfaces...



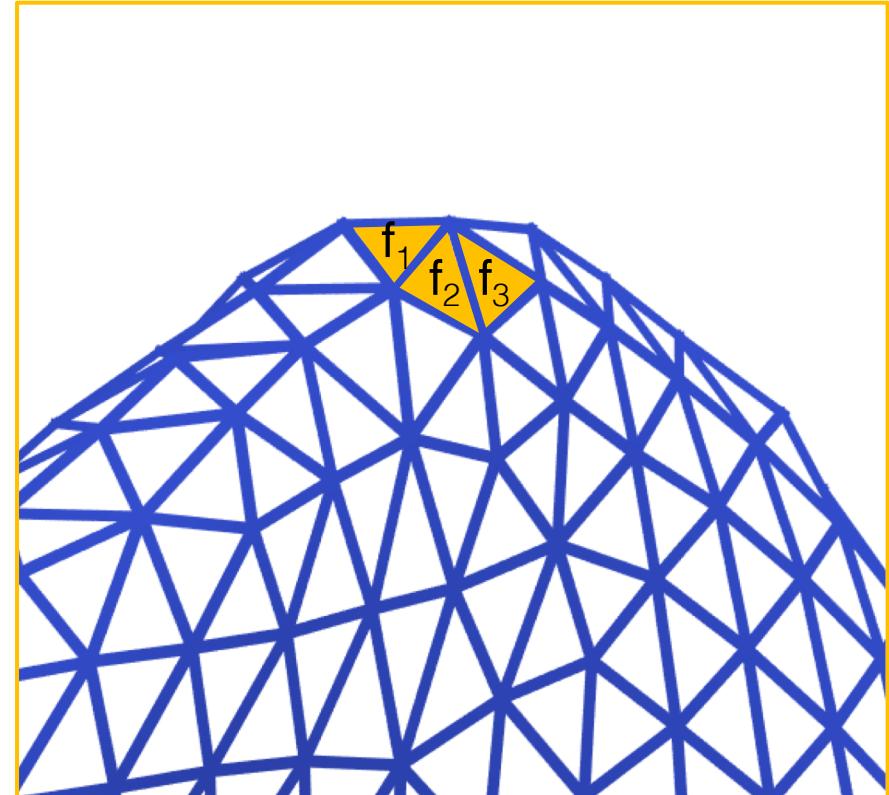
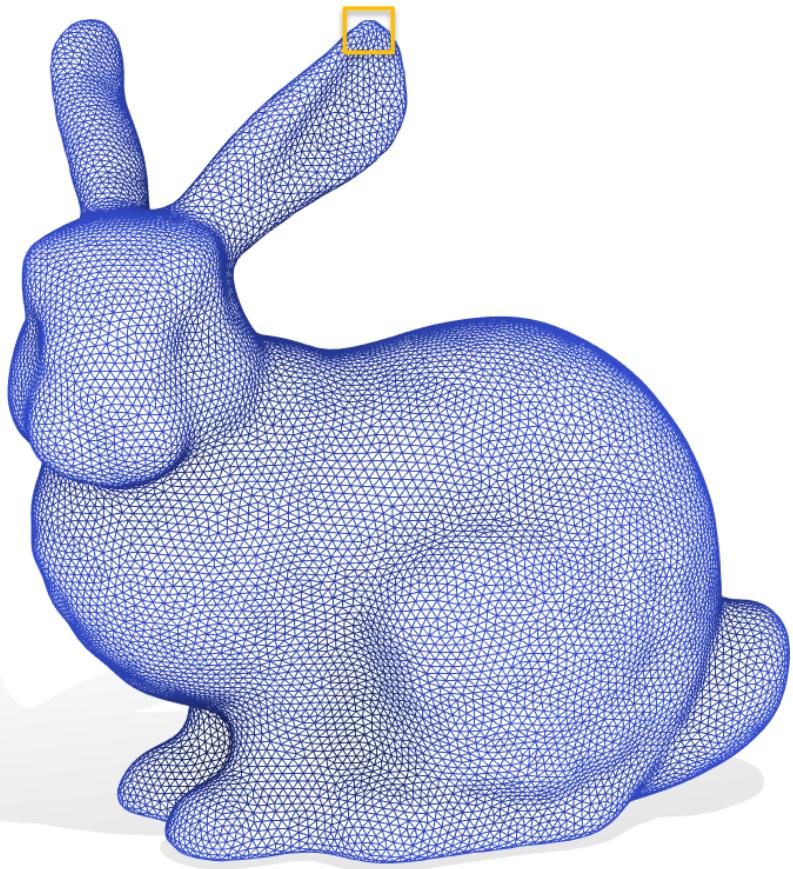
Triangle meshes discretize surfaces...



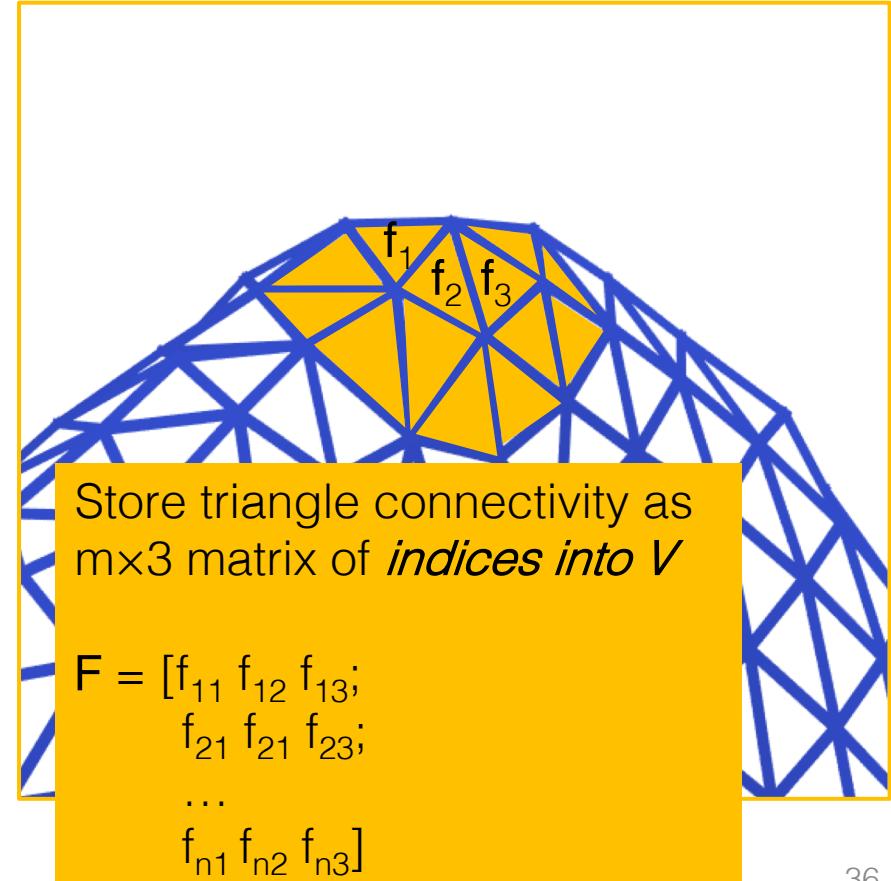
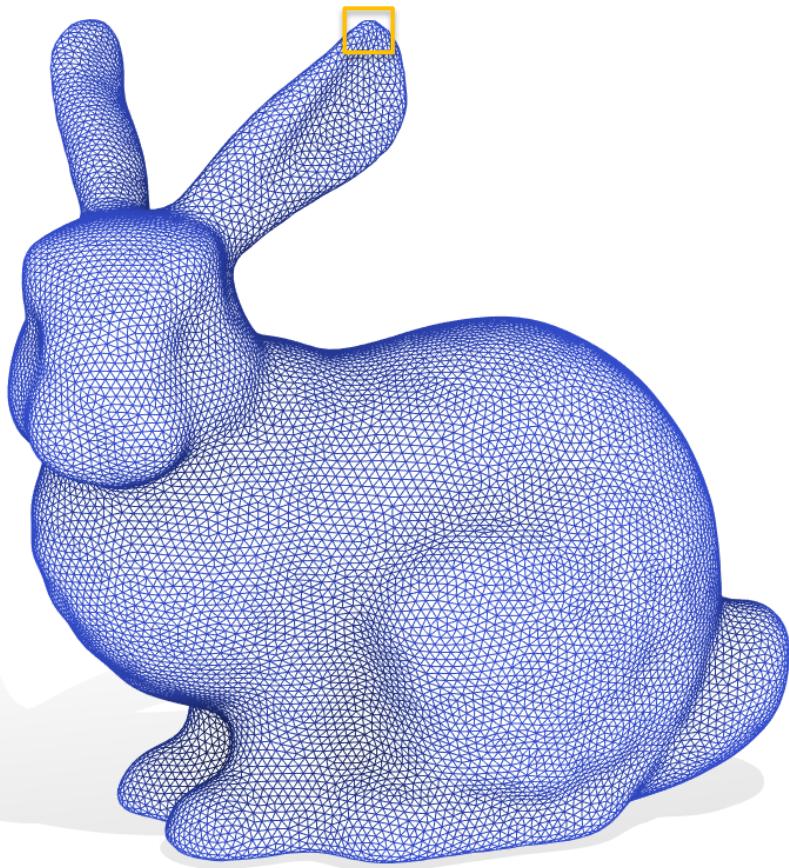
Triangle meshes discretize surfaces...



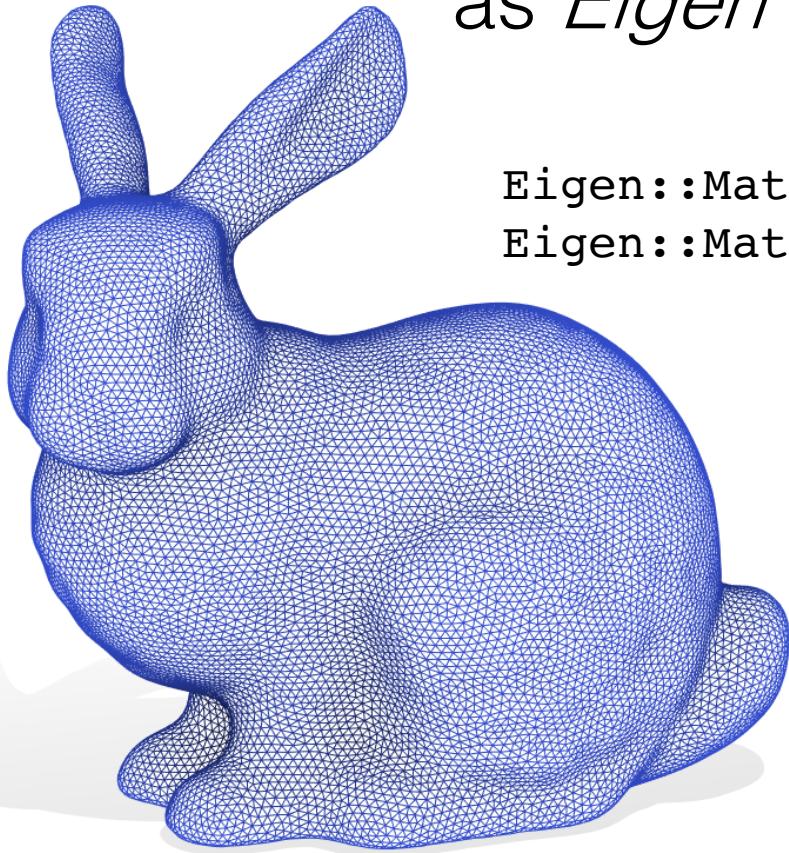
Triangle meshes discretize surfaces...



Triangle meshes discretize surfaces...

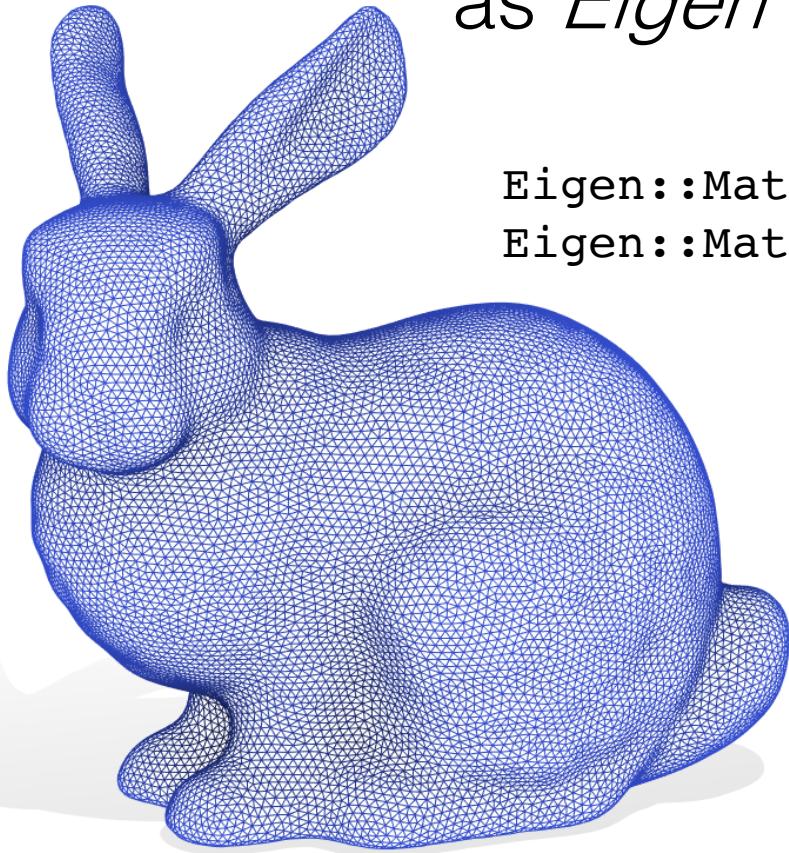


Libigl expects vertex positions and triangle indices
as *Eigen* matrices



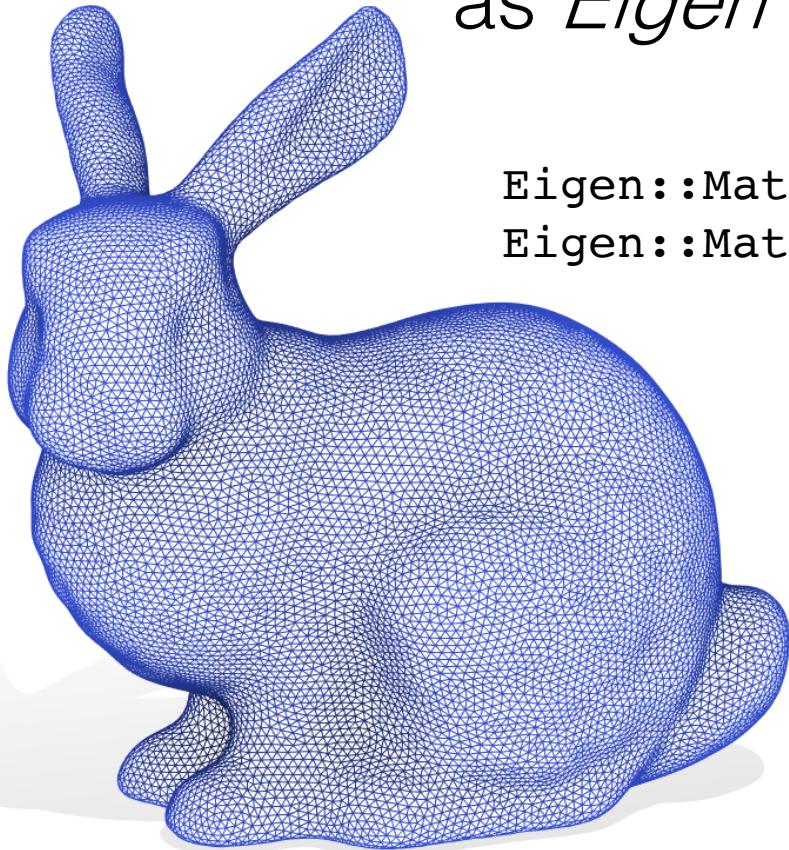
```
Eigen::Matrix<double,Eigen::Dynamic,3> v;  
Eigen::Matrix<int,Eigen::Dynamic,3> F;
```

Libigl expects vertex positions and triangle indices
as *Eigen* matrices



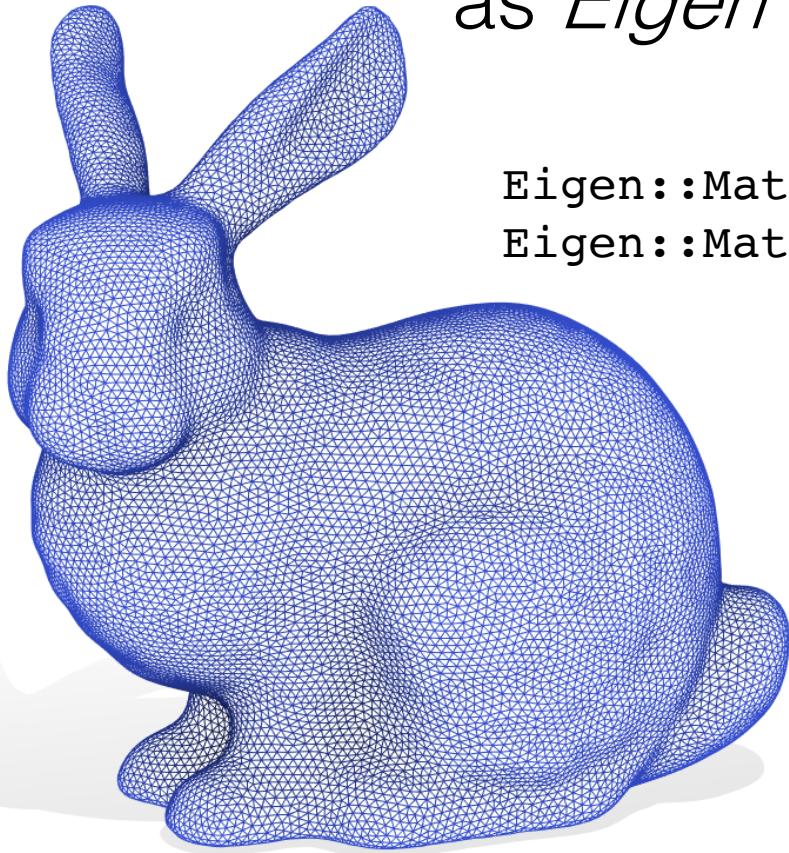
```
Eigen::Matrix<double,Eigen::Dynamic,3> v;  
Eigen::Matrix<int,Eigen::Dynamic,3> F;
```

Libigl expects vertex positions and triangle indices
as *Eigen* matrices



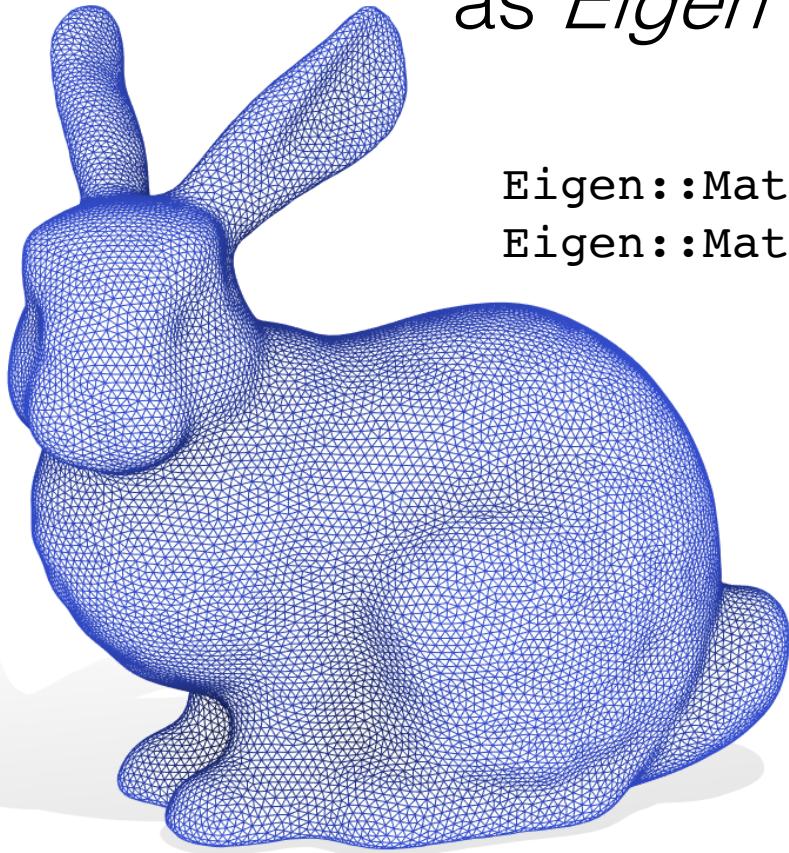
```
Eigen::Matrix<double,Eigen::Dynamic,3> v;  
Eigen::Matrix<int,Eigen::Dynamic,3> F;
```

Libigl expects vertex positions and triangle indices
as *Eigen* matrices



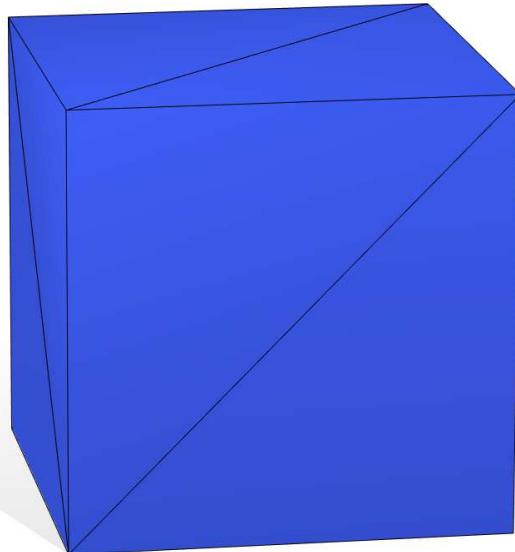
```
Eigen::Matrix<double,Eigen::Dynamic,3> v;  
Eigen::Matrix<int,Eigen::Dynamic,3> F;
```

Libigl expects vertex positions and triangle indices
as *Eigen* matrices



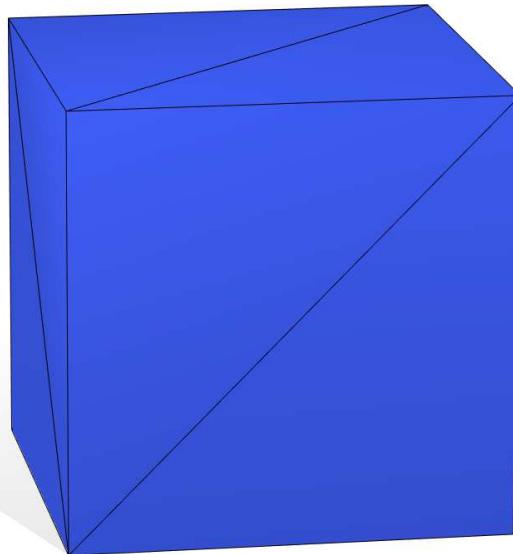
```
Eigen::MatrixXd v;  
Eigen::MatrixXi f;
```

Libigl expects vertex positions and triangle indices as *Eigen* matrices



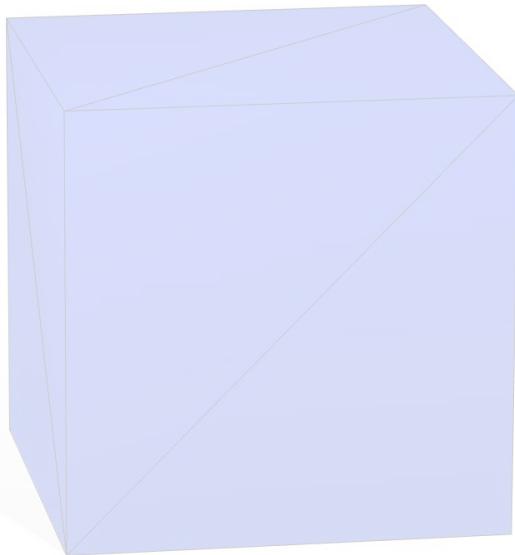
```
Eigen::MatrixXd V(8,3);
V<<
    0.0, 0.0, 0.0,
    0.0, 0.0, 1.0,
    0.0, 1.0, 0.0,
    0.0, 1.0, 1.0,
    1.0, 0.0, 0.0,
    1.0, 0.0, 1.0,
    1.0, 1.0, 0.0,
    1.0, 1.0, 1.0;
Eigen::MatrixXi F(12,3);
F<<
    0, 6, 4,
    0, 2, 6,
    0, 3, 2,
    0, 1, 3,
    2, 7, 6,
    2, 3, 7,
    4, 6, 7,
    4, 7, 5,
    0, 4, 5,
    0, 5, 1,
    1, 5, 7,
    1, 7, 3;
```

Libigl expects vertex positions and triangle indices
as *Eigen* matrices



```
Eigen::MatrixXd V;  
Eigen::MatrixXi F;  
igl::read_triangle_mesh("cube.obj", V, F);
```

Libigl expects vertex positions and triangle indices
as *Eigen* matrices



Eigen:

Eigen:

igl::ref

f 1 3 7

f 1 4 3

f 1 2 4

f 3 8 7

f 3 4 8

f 5 7 8

f 5 8 6

f 1 5 6

f 1 6 2

f 2 6 8

f 2 8 4

cube.obj

v 0 0 0

v 0 0 1

v 0 1 0

v 0 1 1

v 1 0 0

v 1 0 1

v 1 1 0

v;

F;

igl::read_mesh("cube.obj", V, F);

Raw matrices are ...

- memory efficient and cache friendly,
- indices are (often) simpler to debug than pointers,
- trivially copied and serialized,
- immediately passable to other libraries:
OpenGL, OpenCV, MATLAB, embree, ...

Vertex position matrices
immediately *afford* linear algebra

$$\mathbf{u} = \mathbf{v} - \lambda \Delta \mathbf{v}$$



Vertex position matrices
immediately *afford* linear algebra

$$\mathbf{u} = \mathbf{v} - \lambda \Delta \mathbf{v}$$



```
Eigen::MatrixXd U = V - lambda*L*V;
```

Vertex position matrices
immediately *afford* linear algebra

$$\mathbf{u} = \mathbf{v} - \lambda \Delta \mathbf{v}$$



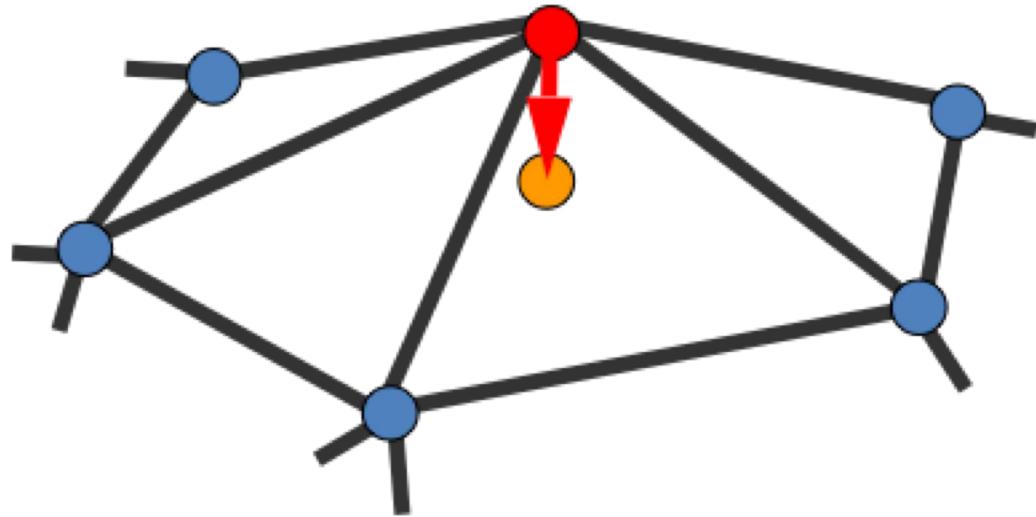
```
Eigen::MatrixXd U = V - lambda*L*V;
```



```
Eigen::SparseMatrix<double> L;
```

Vertex position matrices
immediately *afford* linear algebra

$$\mathbf{u}_i = \mathbf{v}_i - \lambda(\Delta\mathbf{v})_i$$



Vertex position matrices immediately *afford* linear algebra

$$\mathbf{u}_i = \mathbf{v}_i - \lambda(\Delta\mathbf{v})_i$$

$$\mathbf{u}_i = \mathbf{v}_i - \lambda \underbrace{\left(\sum_{j \in N(i)} w_{ij} (\mathbf{v}_i - \mathbf{v}_j) \right)}_{\text{linear in } \mathbf{V}}$$

Vertex position matrices
immediately *afford* linear algebra

$$\mathbf{u}_i = \mathbf{v}_i - \lambda(\Delta\mathbf{v})_i$$

$$\mathbf{u}_i = \mathbf{v}_i - \lambda \underbrace{\left(\sum_{j \in N(i)} w_{ij} (\mathbf{v}_i - \mathbf{v}_j) \right)}_{\text{linear in } \mathbf{V}}$$

$$\mathbf{U} = \mathbf{V} - \lambda \mathbf{L} \mathbf{V}$$

Vertex position matrices immediately *afford* linear algebra

$$\mathbf{u}_i = \mathbf{v}_i - \lambda(\Delta\mathbf{v})_i$$

$$\mathbf{u}_i = \mathbf{v}_i - \lambda \underbrace{\left(\sum_{j \in N(i)} w_{ij} (\mathbf{v}_i - \mathbf{v}_j) \right)}_{\text{linear in } \mathbf{V}}$$

$$\mathbf{U} = \mathbf{V} - \lambda \mathbf{L} \mathbf{V}$$

acts on each column (coordinate) in \mathbf{V}

Vertex position matrices immediately *afford* linear algebra

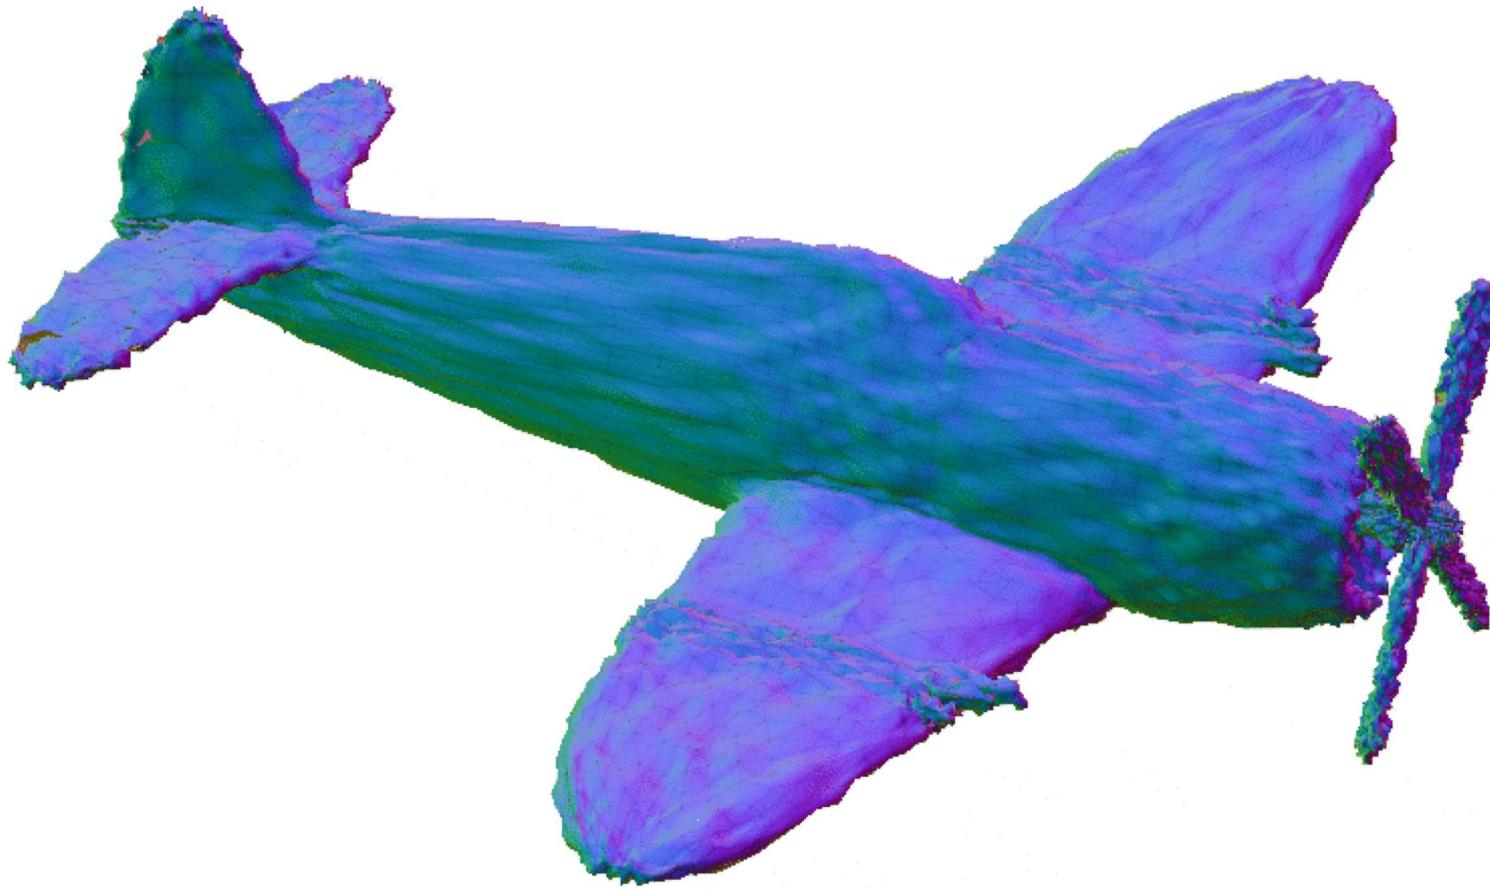
$$\mathbf{u}_i = \mathbf{v}_i - \lambda(\Delta\mathbf{v})_i$$

$$\mathbf{u}_i = \mathbf{v}_i - \lambda \underbrace{\left(\sum_{j \in N(i)} w_{ij} (\mathbf{v}_i - \mathbf{v}_j) \right)}_{\text{linear in } \mathbf{V}}$$

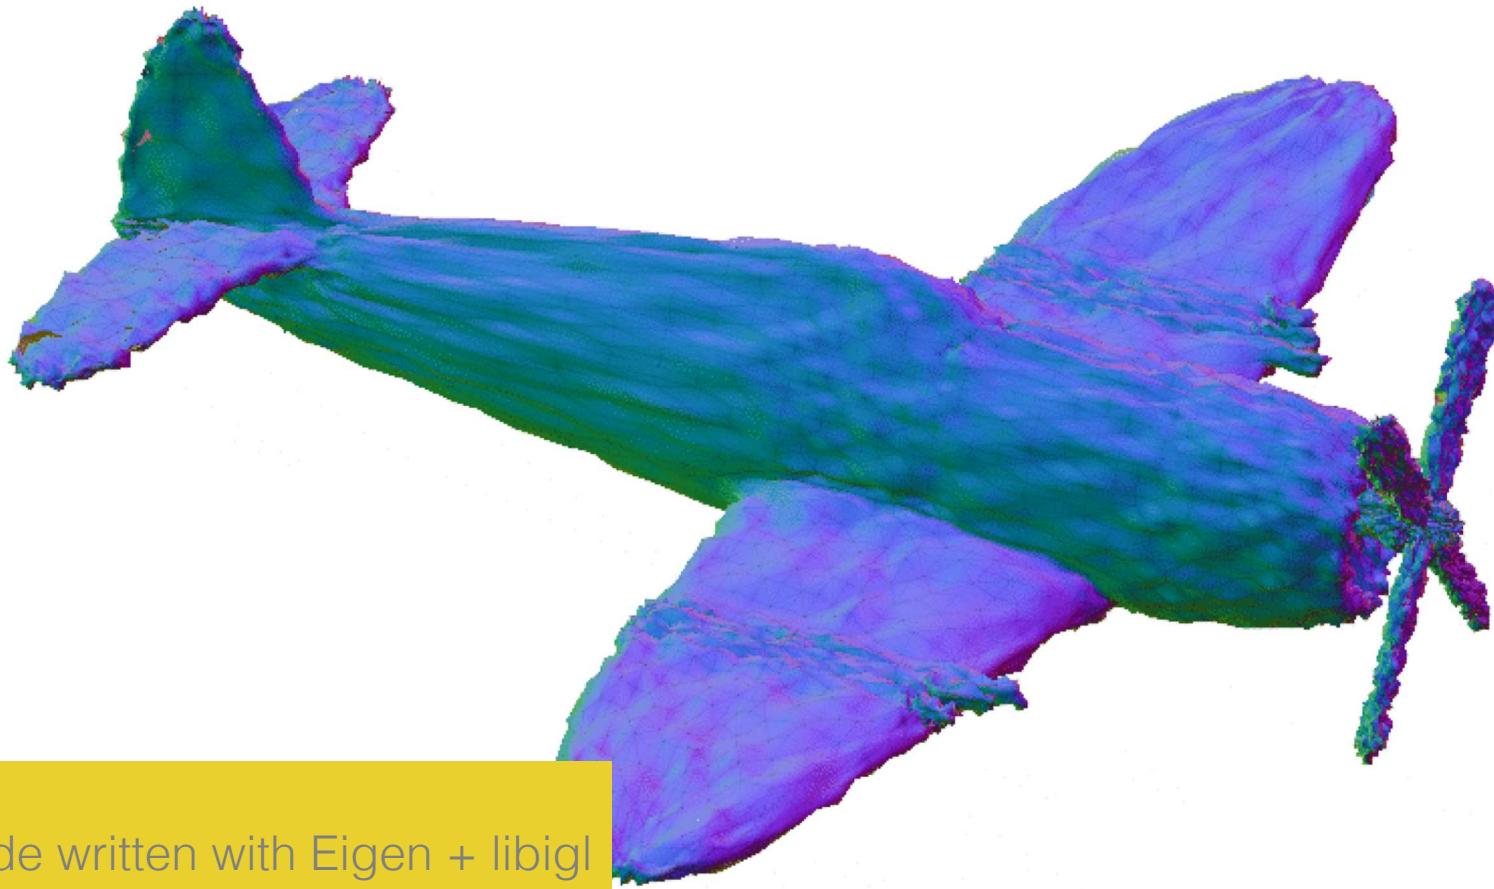
$$\mathbf{U} = \mathbf{V} - \lambda \mathbf{L} \mathbf{V}$$

math \Rightarrow code

```
Eigen::MatrixXd U = V - lambda*L*V;
```



```
Eigen::MatrixXd U = V - lambda*L*V;
```



C++ code written with Eigen + libigl
looks a lot like MATLAB or Python

$$U = V - \lambda \text{ambda} * L * V;$$

Other Matlab-style functions

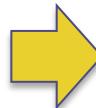
Libigl implements a variety of other routines with the same api and functionality as common Matlab functions.

Name	Description
<code>igl::any_of</code>	Whether any elements are non-zero (true)
<code>igl::cat</code>	Concatenate two matrices (especially useful for dealing with Eigen sparse matrices)
<code>igl::ceil</code>	Round entries up to nearest integer
<code>igl::cumsum</code>	Cumulative sum of matrix elements
<code>igl::colon</code>	Act like Matlab's <code>:</code> , similar to Eigen's <code>Linspaced</code>
<code>igl::cross</code>	Cross product per-row
<code>igl::dot</code>	dot product per-row
<code>igl::find</code>	Find subscripts of non-zero entries
<code>igl::floor</code>	Round entries down to nearest integer
<code>igl::histc</code>	Counting occurrences for building a histogram
<code>igl::hsv_to_rgb</code>	Convert HSV colors to RGB (cf. Matlab's <code>hsv2rgb</code>)
<code>igl::intersect</code>	Set intersection of matrix elements.
<code>igl::isdiag</code>	Determine whether matrix is diagonal
<code>igl::jet</code>	Quantized colors along the rainbow.
<code>igl::median</code>	Compute the median per column
<code>igl::mode</code>	Compute the mode per column
<code>igl::null</code>	Compute the null space basis of a matrix
<code>igl::nchoosek</code>	Compute all k-size combinations of n-long vector
<code>igl::orth</code>	Orthogonalization of a basis
<code>igl::parula</code>	Generate a quantized colormap from blue to yellow
<code>igl::randperm</code>	Generate a random permutation of [0,...,n-1]
<code>igl::rgb_to_hsv</code>	Convert RGB colors to HSV (cf. Matlab's <code>rgb2hsv</code>)
<code>igl::setdiff</code>	Set difference of matrix elements
<code>igl::sort</code>	Sort elements or rows of matrix
<code>igl::speye</code>	Identity as sparse matrix
<code>igl::sum</code>	Sum along columns or rows (of sparse matrix)
<code>igl::unique</code>	Extract unique elements or rows of matrix

libigl embraces matrices as primary mesh data structure

input

```
Eigen::MatrixXd v; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



libigl embraces matrices as primary mesh data structure

input

```
Eigen::MatrixXd v; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



output

```
Eigen::VectorXd x; // #v by 1
```

libigl embraces matrices as primary mesh data structure

input

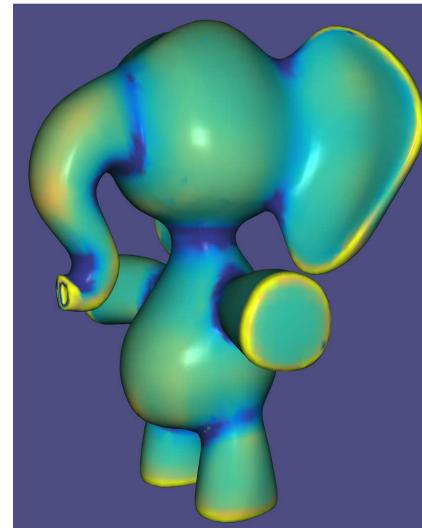
```
Eigen::MatrixXd V; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



output

```
Eigen::VectorXd X; // #v by 1
```

```
igl::gaussian_curvature(V,F,X);
```



libigl embraces matrices as primary mesh data structure

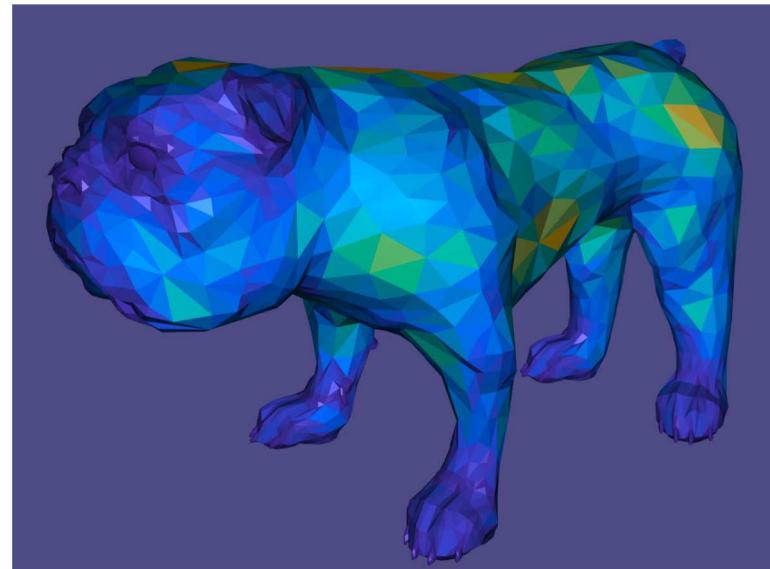
input

```
Eigen::MatrixXd V; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



output

```
Eigen::VectorXd x; // #F by 1
```

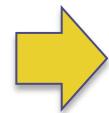


```
igl::doublearea(V,F,X);
```

libigl embraces matrices as primary mesh data structure

input

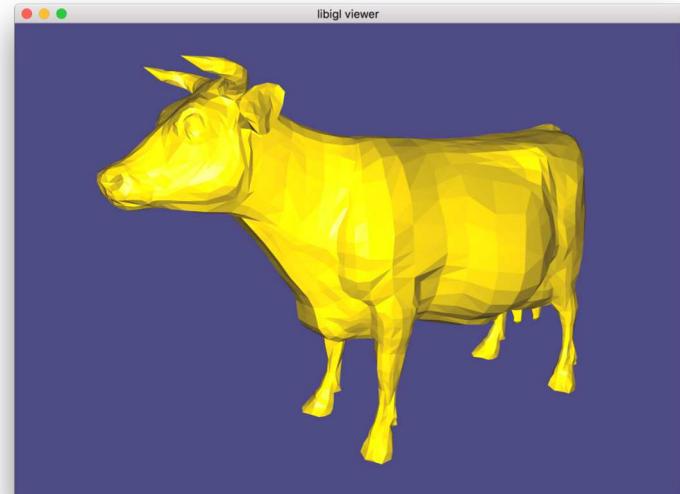
```
Eigen::MatrixXd V; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



output

```
Eigen::MatrixXd X; // #F by 3
```

```
igl::per_face_normals(V,F,X);
```



libigl embraces matrices as primary mesh data structure

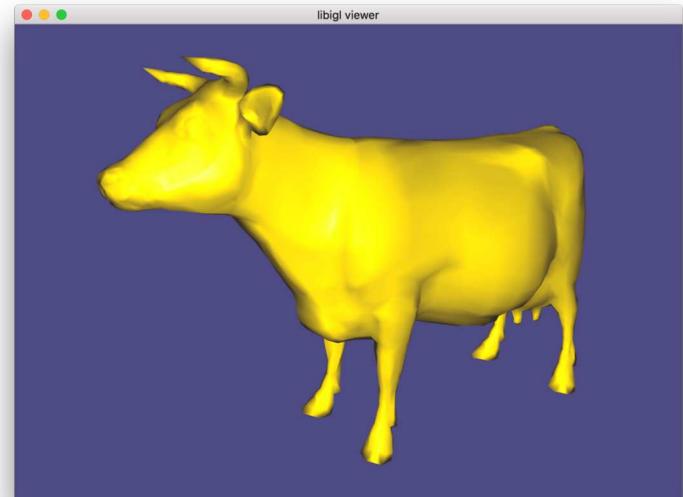
input

```
Eigen::MatrixXd V; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



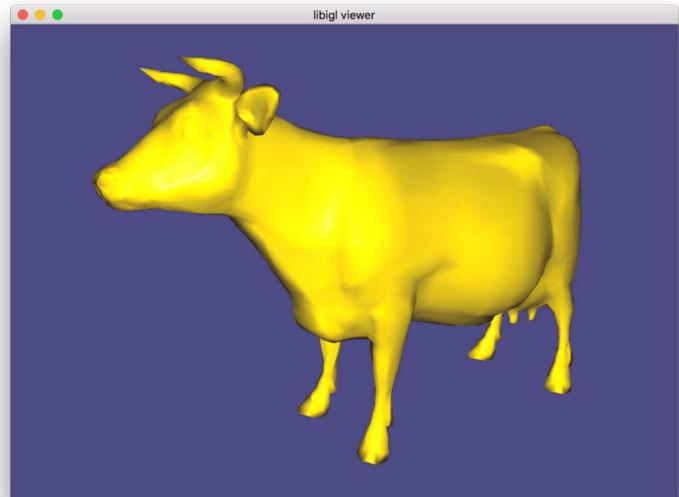
output

```
Eigen::MatrixXd X; // #V by 3
```



```
igl::per_vertex_normals(V,F,X);
```

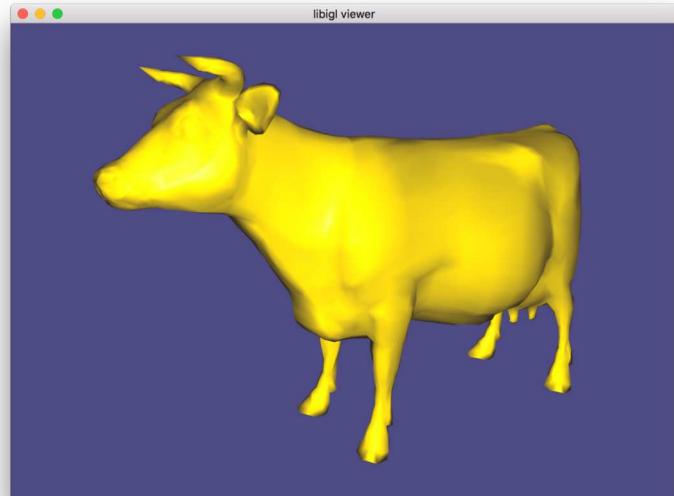
Wait! What about half-edge data-structures?
How do I loop over the one-ring?



Wait! What about half-edge data-structures? How do I loop over the one-ring?

Standard loop over vertices

```
for vertex i  
  
    N(i) = [0,0,0];  
  
    for face f adjacent to i  
  
        N(i) += f's normal  
  
    N(i).normalize
```

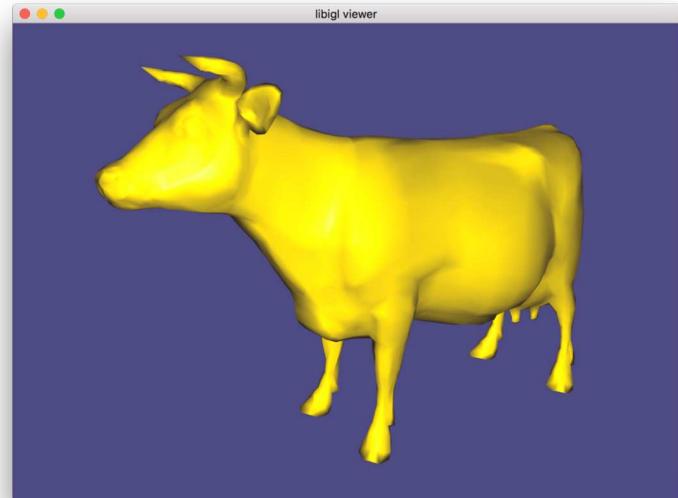


Wait! What about half-edge data-structures? How do I loop over the one-ring?

Standard loop over vertices

```
for vertex i  
  
    N(i) = [0,0,0];  
  
    for face f adjacent to i  
  
        N(i) += f's normal  
  
    N(i).normalize
```

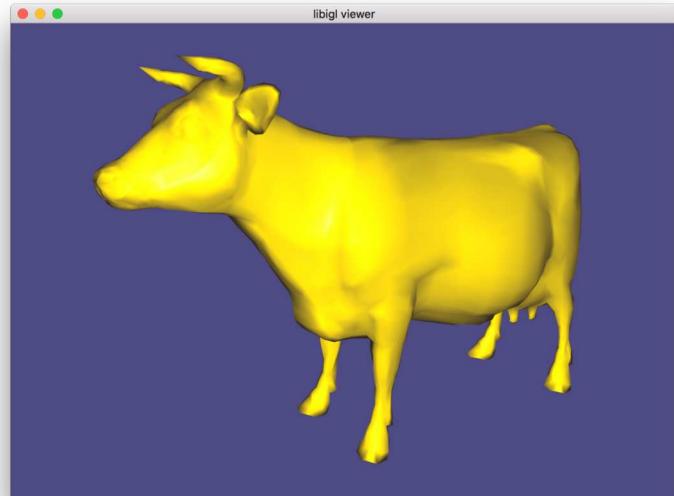
Gather values in place



Wait! What about half-edge data-structures? How do I loop over the one-ring?

Loop over faces

```
N(:,:,:) = 0;  
  
for face f  
  
    for corner vertex i  
  
        N(i) += f's normal  
  
    for vertex i  
  
        N(i).normalize
```

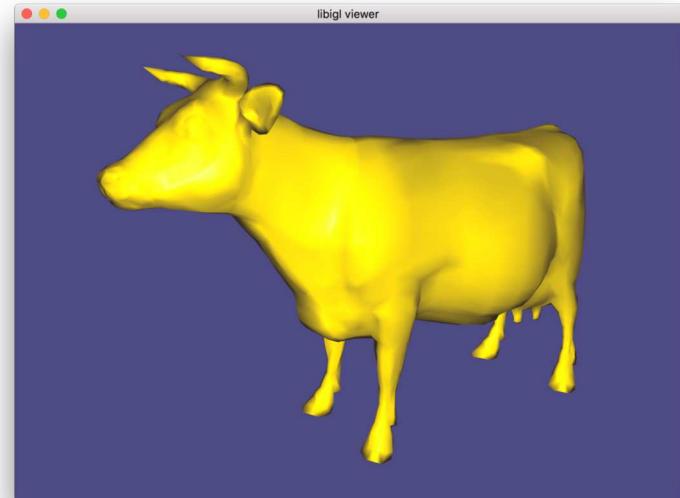


Wait! What about half-edge data-structures? How do I loop over the one-ring?

Loop over faces

```
N(:,:,:) = 0;  
  
for face f  
  
    for corner vertex i  
  
        N(i) += f's normal  
  
    for vertex i  
  
        N(i).normalize
```

Throw values and accumulate



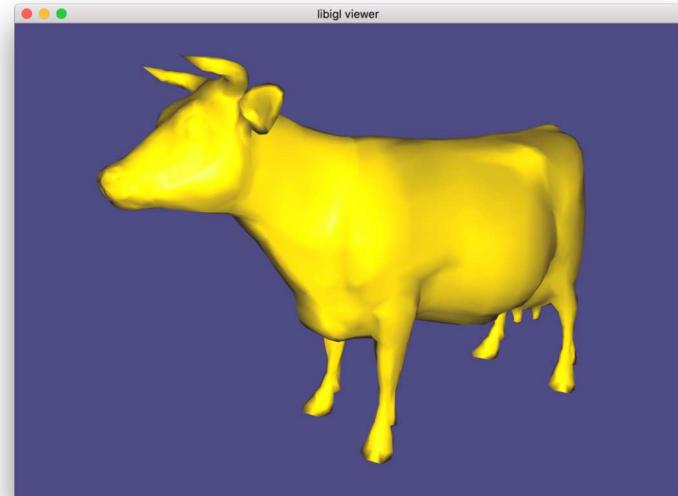
Wait! What about half-edge data-structures? How do I loop over the one-ring?

Loop over faces

```
N = MatrixXd::Zero(n,3);

for(int f=0; f<F.rows(); f++)
    for(int c=0; c<3; c++)
        N.row(F(f,c)) += NF.row(f);

for(int i=0; i<V.rows(); i++)
    N.row(i).normalize();
```



We founded libigl on four core principles

1. Simple, matrix-based API

input

output

```
Eigen::MatrixXd V; // #V by 3  
Eigen::MatrixXi F; // #F by 3
```



```
Eigen::MatrixXd X; // #v by 3
```

```
igl::per_vertex_normals(V,F,X);
```

We founded libigl on four core principles

1. Simple, matrix-based API
2. Header-only “installation”

```
#include <igl/read_triangle_mesh.h>
#include <igl/per_vertex_normals.h>
int main()
{
    igl::read_triangle_mesh("cow.obj",V,F);
    igl::per_vertex_normals(V,F,N);
}
```

We founded libigl on four core principles

1. Simple, matrix-based API
2. Header-only “installation”
3. One function One file

```
// libigl/include/igl/per_vertex_normals.h
#ifndef IGL_PER_VERTEX_H
#define IGL_PER_VERTEX_H
namespace igl
{
    void per_vertex_normals(
        const Eigen::MatrixXd & V,
        const Eigen::MatrixXd & F,
        Eigen::MatrixXd & N);
}
```

We founded libigl on four core principles

1. Simple, matrix-based API
2. Header-only “installation”
3. One function One file
4. Zero dependencies

We founded libigl on four core principles

1. Simple, matrix-based API
2. Header-only “installation”
3. One function One file
4. ~~Zero~~ Minimal dependencies

`libigl/include/igl/*.h`

`e.g., igl::per_vertex_normals()`

depends on stl and Eigen

We founded libigl on four core principles

1. Simple, matrix-based API
2. Header-only “installation”
3. One function One file
4. ~~Zero~~ Minimal dependencies

`libigl/include/igl/*.h`

~70% of libigl

depends on stl and Eigen

We founded libigl on four core principles

1. Simple, matrix-based API
2. Header-only “installation”
3. One function One file
4. ~~Zero~~ Minimal dependencies

`libigl/include/igl/tetgen/* .h`

`e.g., igl::tetgen::tetrahedralize()`

*depends on stl and Eigen, **and Tetgen***

We founded libigl on four core principles

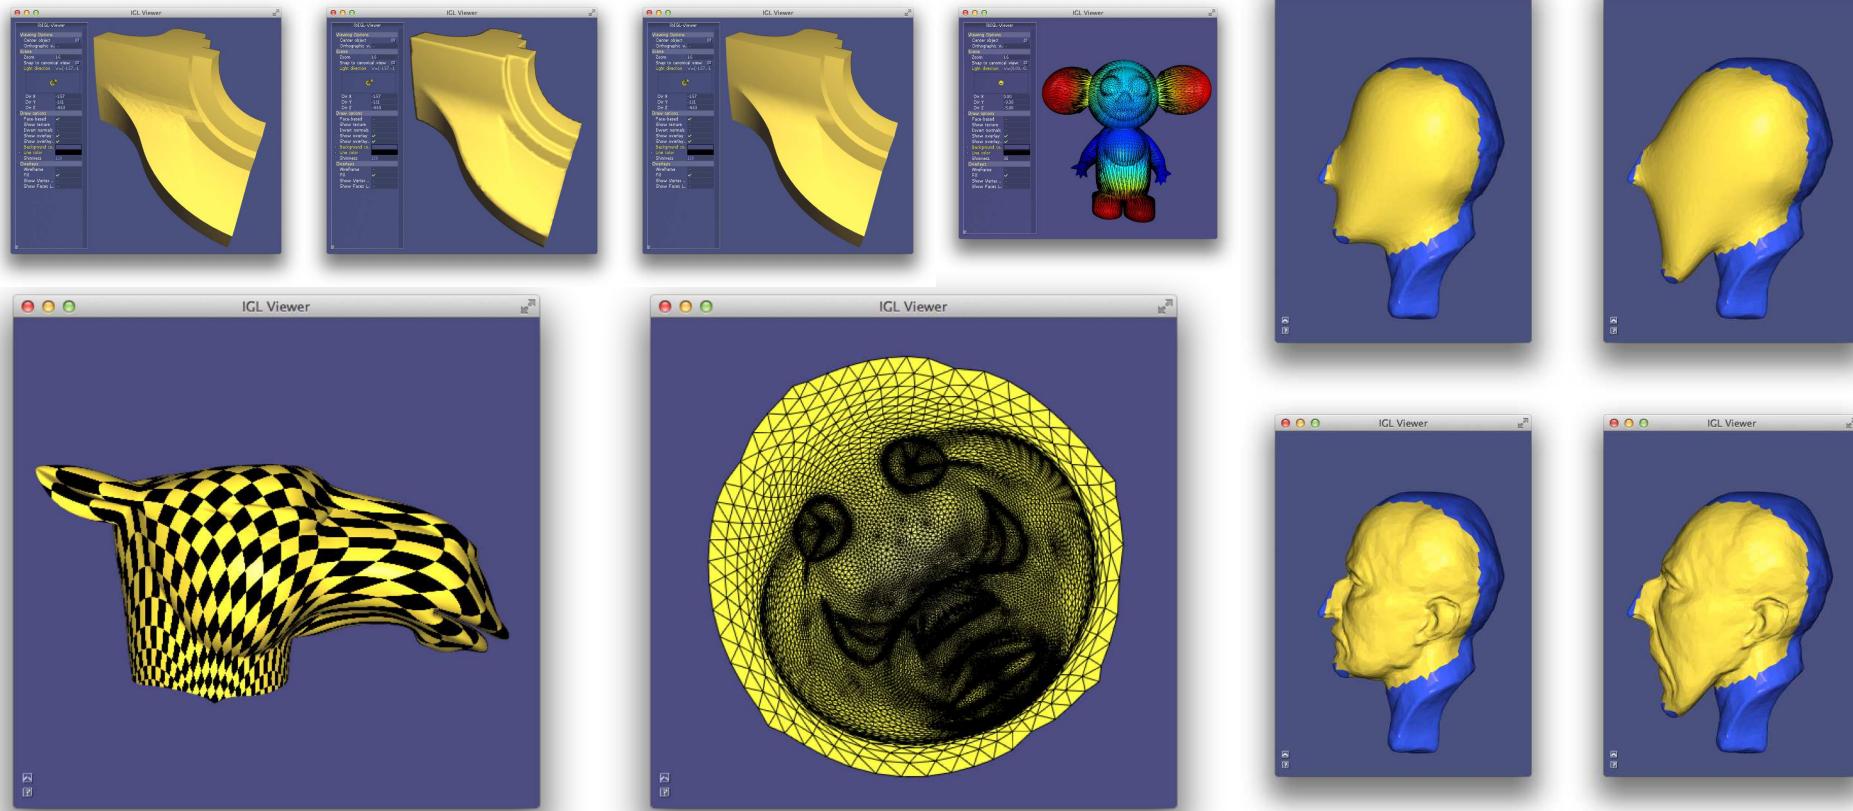
1. Simple, matrix-based API
2. Header-only “installation”
3. One function One file
4. ~~Zero~~ Minimal dependencies

`libigl/include/igl/copyleft/*.h`

`e.g., igl::copyleft::marching_cubes()`

*GPL-type license, **warning** for commerical users*





libigl tour of core functionalities

Tour of Core Functionalities

- <https://github.com/libigl/libigl-course>
 - Tutorials compilation with CMAKE
 - Differential Quantities
 - Parameterization
 - Deformation (and Locally Injective Maps)

Get started with:

```
git clone --recursive https://github.com/libigl/libigl.git
```

libigl Getting Up and Running with libigl

Getting Up and Running with libigl

- The core of libigl is header only
- The best way to start a new project is to use the basic CMAKE project
- Useful convenience tools
 - Render on memory buffer (screenshots)
 - Picking
 - Menu bar (nanogui)

libigl hot new features

New Features

- Serialization
- Booleans
- Python

Libigl Coding Tips (aka “How to code a SIGGRAPH project”)

This is a short list of coding tips that will greatly reduce your pain and suffering before (and after) the SIGGRAPH deadline.

1. Serialize it all

The entire state of your application should be serializable, i.e. It should be possible to save it into a binary file and reload it at any point. This drastically simplifies debugging, since you can serialize just before a crash happens and debug from that point without running your complete algorithm again. Serializing all results shown in the paper’s figures enables quicker editing iterations before (and after) the deadline. It also allows you to share your results with others that wish to compare with your method. An additional tip is to serialize the state of the application on the window close event and automatically reload it when you launch it again.

2. Always assert

libigl coding tips

A screenshot of a GitHub repository page. The repository is named `alecjacobson / geometry-processing`. The page shows basic statistics: 1 commit, 1 branch, 0 releases, and 1 contributor. There are buttons for creating a new file, uploading files, finding files, and cloning or downloading the repository. The GitHub interface includes a header with navigation icons, a search bar, and links for pull requests, issues, marketplace, and gist.

Course material for a grad-level course in Geometry Processing.

geometry-processing libigl Manage topics

1 commit 1 branch 0 releases 1 contributor

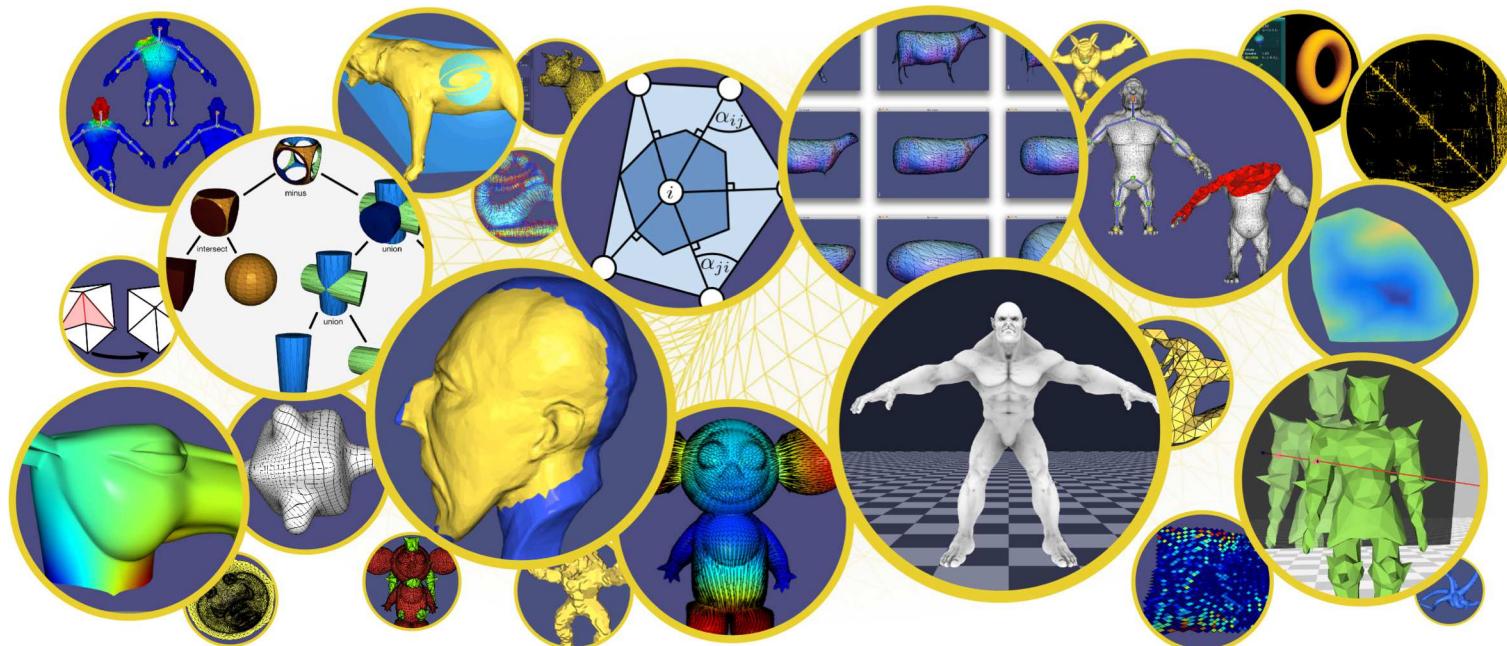
Branch: master New pull request Create new file Upload files Find file Clone or download

libigl open-source grad course in geometry processing

Another geometry processing course using libigl

<https://github.com/danielepanozzo/gp>

libigl Prototyping Geometry Processing Research in C++



Alec Jacobson
University of Toronto

<https://github.com/libigl/libigl-course>

Daniele Panozzo
New York University